

Goals-First Kustomer (Name tbd)

A vision for how Kustomer wins the next era of customer experience

Author	Aditya Subramaniam
Status	Draft for executive review
Primary audience	Kustomer executive team
Secondary audience	Broader Kustomer org
Date	May 2026

Executive Summary

The opportunity

Every CX organization has business goals: deflect more tickets, improve CSAT, reduce time-to-resolution, save more at-risk accounts, drive more expansion revenue. Today, no platform in the market is built around them.

Every competitor sells the *means*: AI agents, copilots, dashboards, QA tools, automation builders. Admins have to assemble those means into something that drives their goals. They configure procedures, build scorecards, set thresholds, wire up reports, and then hope it all adds up. When it doesn't, they have no systematic way to know why, or what to change.

Intercom, Zendesk, Sierra, and Decagon have all gestured at this gap. Intercom's CEO promised "Goals" on stage at Pioneer 2025 ("Fin will pursue your objectives and priorities") — but eight months later it remains future-tense. Sierra talks about "outcomes you define and agents that deliver them," but ships per-agent guardrails, not workspace-level goals. Zendesk and Decagon use outcome-based *pricing* as a proxy for outcome-based *configuration*. Pylon owns the B2B account-intelligence wedge but has no goals primitive at all.

Goals-first configuration is genuinely open ground. Whoever ships it first owns the category vocabulary, the analyst narrative, and the next three years of category-defining product marketing.

The vision: Goals-First Kustomer

Kustomer becomes the definitive platform that ensures CX organizations achieve their business goals using AI and humans.

Every part of Kustomer aligns to the goals an admin declares:

- **Admin setup** asks first: "What outcomes do you need to achieve?" Goals drive the entire configuration experience, not the other way around.
- **AI for Customers (AIC)** configures itself against those goals. Procedures, escalation thresholds, guardrails, and tool access are all goal-aware. AIC autonomously resolves what it can, in the way that drives the goals.
- **AI for Reps (AIR)** partners with reps on the escalations. Suggestions, drafted responses, and next-best-actions are tuned to the goals that matter for each conversation and account.
- **AI for Leaders** (data explorer, workflow assistant, performance analysis) reasons in the language of goals. When a leader asks "how are we doing?", the answer is framed against the goals they declared.
- **Reporting and accountability** are organized around goals, not channels or queues. Every conversation, every AI action, every rep handle contributes to (or detracts from) measurable goals.

This is not a feature. This is a re-foundation of how a CX platform is structured. Kustomer becomes the only platform where the data model, the AI, the reps' tools, and the reporting are all unified by a single primitive: **the business goal**.

Why now

Three forces converge in 2026:

1. **Outcome-based pricing is becoming table stakes** (\$0.99/resolution at Intercom, outcome-only contracts at Sierra and Decagon). The market is conditioned for outcome thinking. No one has converted that conditioning into outcome-based *configuration*.
2. **AI governance is becoming the executive concern**. Gartner: "By 2030, 50% of AI agent deployment failures will be due to insufficient AI governance platform runtime enforcement." 91% of CX leaders are under pressure to implement AI in 2026. They need a way to control AI that is expressed in business terms.
3. **Buyers are tired of assembling platforms**. Futurum's 1H 2026 survey found direct revenue and profit impact overtook productivity gains as the primary AI ROI metric. CFOs want goal-aligned spend, not feature-aligned spend.

The window closes in 12 to 18 months. Intercom is closest to converting their stated vision into product. Sierra's Ghostwriter (March 2026) is two months old and aimed at mid-market. Pylon is raising again. Kustomer can win this, but only by moving decisively now.

Beacon: the first milestone

Beacon (name tbd) is the goals-first governance layer. It is the first piece of the vision because it is the one that unlocks everything else:

- **Goals** as a first-class object with thresholds, owners, and trajectories.
- **Monitors (name tbd)** that score every AI and human conversation against those goals.
- **AI Computed Fields (name tbd)** (AI-Generated CSAT, Customer Health Score, Company Health Score) that feed goal progress.
- **Suggestions** that close the loop: when a goal is at risk, Beacon proposes concrete changes.
- **Anomaly Detection** from the AIC 2.0 Validation Agent: runtime drift detection that no competitor ships today.

Beacon ships in milestones starting June 2026. It is the foundation. Everything that follows — goal-aware admin setup, goal-aware AIC configuration, goal-aware AIR, goal-aware AI for Leaders — extends from the primitives Beacon establishes.

What it takes

This vision requires three things from the broader Kustomer org:

1. **Conviction on the primitive.** Goals are not a feature we bolt on. They are the organizing object for the platform. Every product team needs to make their surface goal-aware over the next 6 to 12 months.
2. **Sequencing discipline.** Beacon establishes the primitive. Subsequent milestones build on it. We do not parallel-track this. We sequence to avoid five teams shipping five incompatible definitions of "goal."
3. **A clear external narrative.** "Goals-First Kustomer" must show up in analyst briefings, sales decks, customer kickoffs, and the product itself within the next two quarters, before Intercom converts their roadmap statement into a shipped product.

Mock Press Release

The following is a hypothetical press release written as if Goals-First Kustomer were generally available today. It is included to make the vision concrete for internal alignment, not for external use.

FOR IMMEDIATE RELEASE

Kustomer Launches the First Goals-First CX Platform, Aligning AI and Humans to the Business Outcomes That Matter

New platform architecture puts business goals at the center of every workflow, AI agent, and reporting surface, eliminating the gap between what CX teams aim for and what their software does

NEW YORK, [Date] — Kustomer today announced the launch of the first goals-first customer experience platform, fundamentally redesigning how CX organizations configure their software, deploy AI, and measure success. Rather than asking customers to assemble disconnected tools to chase outcomes, Kustomer now starts every workflow with a single question: "What business goals are you trying to achieve?"

For the first time, a CX platform's administrative setup, AI agents, rep tools, and analytics all share a common primitive: the goal. When a CX leader declares a target — improve auto-resolution rate, reduce time-to-resolve for VIP accounts, save more at-risk renewals — Kustomer aligns the entire platform to drive it.

"For twenty years, CX software has been built around channels, queues, and tickets," said Brad Birnbaum, Co-Founder and CEO of Kustomer. "But none of those things are what businesses actually care about. Businesses care about outcomes. Our customers told us their AI agents were getting smarter but their ability to steer those agents toward business results wasn't keeping up. Goals-First Kustomer fixes that. You tell us what you want to achieve. We configure the platform — AI, humans, and reporting — to deliver it."

How it works

Goals-First Kustomer transforms every part of the platform:

- **Admin setup is goal-led.** Configuring Kustomer starts with declaring the goals the organization needs to achieve. Setup wizards, default configurations, and best-practice templates are all selected to drive those goals.
- **AI for Customers (AIC) is goal-aware.** Kustomer's autonomous AI agent configures its procedures, escalation thresholds, tool access, and guardrails based on the declared goals. Customers see resolutions tuned to what the business actually values.
- **AI for Reps (AIR) partners with humans on the goals.** When a conversation escalates, AIR knows which goals are in play, drafts responses optimized for them, and flags risks before they become regressions.
- **AI for Leaders speaks the language of goals.** Data Explorer and Workflow Assistant answer questions framed against the goals leaders declared, making it possible to ask "how are we doing on renewal saves this quarter?" and get an answer grounded in real conversation data.

- **Beacon governs the loop.** Kustomer's Beacon product continuously evaluates every AI and human conversation against the goals, surfaces concrete suggestions when goals are at risk, and detects runtime anomalies in AI behavior before they impact customers.

Available today

Kustomer's goals-first capabilities are generally available to all customers starting today. Beacon, the first goals-aware product, includes Goals, Quality Monitors, AI Computed Fields (AI-Generated CSAT, Customer Health Score, Company Health Score), Suggestions, Anomaly Detection, and Goal-Level Performance Reporting. Additional goal-aware capabilities across AIC, AIR, and AI for Leaders ship over the coming quarters.

About Kustomer

Kustomer is the AI-native CX platform built for B2C and B2B organizations that need to align AI and human work to business outcomes. Founded in 2015 and headquartered in New York, Kustomer serves leading brands across retail, DTC, B2B SaaS, financial services, and travel.

Media contact: [TBD]

Proposed Roadmap

Sequencing and rapid iteration is what makes this vision deliverable. The principle: **establish the goals primitive once, then make every surface goal-aware**. If five product teams each invent their own definition of "goal," the vision fails. If we sequence correctly, every subsequent product surface gets goals for free.

This roadmap focuses on the steps required, not calendar dates. Dates are owned by individual product teams.

Milestone 1: Land Beacon M1 with Goals as a first-class platform object

Beacon M1 is already scoped (target: Summer 2026). The critical thing for the vision is that Beacon ships **Goals as a first-class platform object**, not as a feature internal to monitoring. Specifically:

- Goals must be addressable from any other Kustomer product surface via API.
- Goals must have stable schemas (field references, thresholds, directions, owners) that AIC, AIR, and AI for Leaders can read.
- The Goals data model must support both B2C goals (CSAT, deflection, AHT) and B2B goals (account health, renewal risk, NRR).

- The Goals UI in Beacon becomes the canonical place to declare and edit goals; every other surface that interacts with goals navigates back to it.

This is the foundation. Without it, the rest of the vision is incoherent.

Milestone 2: Make Kustomer AI goal-aware

Kustomer AI (AIC and AIR together) is the highest-impact surface for goal awareness. The principle is the same across both: the admin declares a goal, and Kustomer AI configures itself to drive it, whether the work is autonomous (AIC) or assisted (AIR).

Steps for AIC (autonomous):

- AIC procedures know which goals they support; the procedure builder surfaces "this procedure contributes to Goal X" inline.
- Escalation thresholds become goal-aware (more aggressive escalation when goal trajectory is at risk).
- The Validation Agent uses goals as part of its drift detection logic: anomalies are scored against goal impact.
- AIC's testing surface (already shipped) reports outcomes in terms of declared goals, not raw metrics.

Steps for AIR (assisted):

- AIR knows which goals are in play for each conversation (based on customer segment, conversation topic, account tier).
- AIR's drafted responses and next-best-actions are scored against goal contributions.
- Rep performance reporting is reframed around goal contribution (not just AHT and CSAT).

Steps for the handoff between them:

- When AIC escalates to AIR, the goals in play travel with the conversation so the rep inherits the same context the AI was optimizing for.
- The same goal-contribution scoring applies whether a conversation was resolved by AIC, by a rep working with the assistance of AIR, or by a rep working without assistance.

Milestone 3: Make admin setup goal-led

Once Goals exist as a primitive, the admin setup experience is rebuilt around them. The first time an admin sets up Kustomer (or revisits config), they declare what they are trying to achieve. The platform then proposes configurations – queues, routing, AIC procedures, AIR setup, KB structure – tuned to those goals.

This is the highest-leverage early bet because it reframes the entire onboarding experience around outcomes. It is also the most visible to prospects in sales motions.

Steps:

- Identify the 5 to 7 most common goals across the customer base (working with CS and the AIR/AIC PMs).
- Build a goal-led setup wizard that maps each goal to a starting configuration template.
- Refit existing setup screens so they show "this configuration contributes to your declared goals" inline.

Milestone 4: Make AI for Leaders goal-aware

Data Explorer and Workflow Assistant are the leader-facing surfaces. Today, leaders ask raw analytical questions. In the goal-aware future, leaders ask questions framed against the goals they declared, and the AI assistant frames its answers the same way.

Steps:

- Data Explorer indexes goal definitions and progress as part of its grounding context.
- Workflow Assistant proposes workflow changes in terms of goal contribution.
- A goal-level reporting surface (extending Beacon's M1 reporting) becomes the default leader landing page.

What success looks like at end of Q3

- Beacon M1 GA with Goals as an addressable platform object.
- Admin setup wizard goal-led, in beta.
- AIC procedure builder shows goal contribution inline.
- Three reference customers running on declared goals.

What success looks like at end of Q4

- AIC, AIR, and AI for Leaders all goal-aware.

- Goal-led admin setup is the default for new orgs.
 - "Goals-First Kustomer" is the primary external narrative.
-

Appendix A: Direct Competitive Analysis

Intercom Fin

- **Shipped (May 2026)**: Fin 3 (66% average resolution; approximately \$100M+ ARR for Fin alone; 8,000 businesses); Procedures (branching workflow automation, 1.5M+ conversations); Simulations (pre-deploy testing with auto-generated personas); Monitors + Custom Scorecards (March 25, 2026); CX Score + CX Score Reasons; Topics Explorer; Recommendations; Trends; Fin Voice (28 languages); Fin for Sales (April 2026); Fin for Ecommerce (May 8, 2026); Slack + Discord channels.
- **Goals orientation: Vision only**. Pioneer 2025 (October 9, 2025) announced Roles, Goals, Memory, Knowledge, Interoperability as future capabilities. Goals NOT shipped as of May 2026. Closest analog is the Fin for Sales "Playbook," a natural-language instructions doc; role-scoped, not workspace.
- **QA rigor**: Strong. Monitors + Custom Scorecards is among the most mature shipped AI-eval primitives. Procedure Monitors and Knowledge Base Checks on roadmap.
- **Pricing**: Fin \$0.99/resolution; helpdesk from \$29/seat/month; Copilot \$35/month; **Pro add-on \$99/month paywalls Monitors, Custom Scorecards, CX Score, Topics, Trends, Recommendations** (effective March 12, 2026 for new customers; May 12, 2026 cutoff for legacy users). Net revenue retention (NRR) for new customers reportedly went from 112% to 146% on Fin outcome pricing (CEO Eoghan McCabe board update, June 2025).
- **B2C vs B2B**: B2B-leaning, with strong mid-market and SMB penetration.
- **Threat level: Highest**. Strongest combination of shipped capability, brand mindshare, momentum, and stated goals vision.
- **Customer wedge**: Pro paywall creates pricing friction at \$99/month. Beacon bundles Monitors, Scorecards, AI Computed Fields into the core platform. Also: Intercom's Goals is still vision; Kustomer can ship the primitive first.

Zendesk Resolution Platform

- **Shipped**: AI Agents (next-gen with adaptive reasoning); Knowledge Graph + Knowledge Builder; Action Builder + App Builder; AI Reasoning Controls (chain-of-thought audit); Custom QA (100% scoring of human + AI); AI Insights Hub; Generative Search; Voice via Local Measure; Copilot enhancements with autonomous procedure execution.

- **Forethought acquisition:** Announced March 11, 2026, **completed March 26, 2026**. Brings "Resolution Learning Loop" (detects workflow gaps, generates new procedures, tests optimizations). Strong Performer in Forrester Wave Customer Service Solutions Q1 2026.
- **Goals orientation:** Outcome-based **pricing** and "resolution" identity, but no admin-declared goals primitive.
- **QA rigor:** Strong via Custom QA + AutoQA + Reasoning Controls.
- **Pricing:** Outcome-based for AI Agents; per-seat for core helpdesk.
- **B2C vs B2B:** Both, with explicit "Specialized AI agents for B2B, B2C, and B2E" framing post-Forethought.
- **Threat level: High.** Distribution moat plus Forethought puts them within 6 to 12 months of feature parity on the QA/governance loop.
- **Customer wedge:** Zendesk is built around channels and tickets; Kustomer can be built around goals. Also: Zendesk's data model is fragmented across acquisitions; Kustomer's unified CRM is structurally simpler for goal-aware configuration.

Sierra (Agent OS)

- **Shipped:** Agent OS (chat, SMS, WhatsApp, email, voice, native ChatGPT integration); Agent Data Platform (November 2025, unifies unstructured + structured customer data with "strategies" defined by audience, outcomes, inventory, triggers); **Ghostwriter** (March 25, 2026 — agent-builds-agents, 30+ languages, SOPs/transcripts/audio inputs); Explorer (agent-optimizing agent); Simulations; voice agents at sub-1-second latency.
- **Goals orientation:** Closest competitor philosophically. Bret Taylor: "just outcomes you define and agents that deliver them." Reality: per-agent goals and guardrails, not workspace-level.
- **Mid-market push:** Ghostwriter explicitly designed to extend Sierra into mid-market that lacks engineering teams. Until now Sierra was 4 to 10 weeks sales-led with forward-deployed engineers; entry-level contracts approximately \$150K/year (SelectHub, third-party); Year-1 all-in \$200K to \$350K+ with setup. Ghostwriter is two months old at the time of writing; no public validation that mid-market self-serve is working.
- **Pricing:** Outcome-only. Charges when AI resolves an interaction, saves a cancellation, or completes an upsell. No per-seat or per-usage.
- **Recent funding:** Series E raised \$950M at \$15.8B post-money valuation (Tiger Global + GV leads); \$150M ARR as of February 2026.
- **B2C vs B2B:** Both. Heavy regulated industries (Cigna, Prudential, Rocket Mortgage, Blue Cross Blue Shield, 1 in 3 of the world's largest banks) plus DTC retail (Nordstrom, Vivid Seats).
- **Threat level: High at enterprise.** Less direct B2C mid-market threat than Intercom; more enterprise-vertical.
- **Customer wedge:** Sierra's \$150K ACV floor limits mid-market reach. Kustomer covers AI + human + reporting in one platform; Sierra is AI-only.

Decagon

- **Shipped:** Agent Operating Procedures (AOPs); **Duet** (March 2026, AI partner that runs the agent lifecycle via conversation on top of AOPs); Watchtower (always-on QA, every conversation); Simulations (test against historical transcripts + AI-generated personas); Experiments (live A/B traffic split, CSAT/deflection-tracked); Knowledge Suggestions; AI Actions (Stripe, Shopify, Salesforce); Proactive Agents (user memory + outbound voice); Agent Workbench (autonomous debugging); intent tags surfacing customer goals; Git-based agent versioning.
- **Goals orientation:** Strong outcomes language and "Agent Engineers" framing; per-AOP goals, but no workspace-level goals primitive.
- **QA rigor:** Among the strongest. Watchtower + Simulations + Experiments is the most rigorous shipped eval trio after Sierra.
- **Pricing:** Enterprise-only; no public pricing; no self-serve.
- **B2C vs B2B:** Both; heavy mid-market and enterprise B2C (subscription, fintech).
- **Threat level: Medium-high.** Best-in-class agent engineering depth; less brand reach.
- **Kustomer wedge:** Decagon is agent-engineering-led (sells to engineering teams). Kustomer is goal-led (sells to business users).

DigitalGenius

- **Shipped:** Care AI, Purchase AI, Proactive AI, Voice AI, Insights AI; Genius Flows (pre-built ecommerce workflows for WISMO, returns, warranty, subscription management); Visual AI (product defect detection from images, barcode reading); deep integrations with Shopify, Adyen, FedEx, DHL, eBay; integrates with Kustomer, Zendesk, Salesforce, Gorgias.
- **Goals orientation:** None at the product level. Marketing positions around "ecommerce outcomes" but no goals primitive.
- **QA rigor:** Light; relies on customer success team for active management.
- **Pricing:** Custom; sales-led.
- **B2C vs B2B:** Pure-play ecommerce/B2C.
- **Threat level: Low-medium.** Strong in fashion/DTC mid-market; not a horizontal CX platform.
- **Kustomer wedge:** DigitalGenius is narrow to fashion/DTC; Kustomer is horizontal with goals as the unifying primitive.

Pylon (B2B-focused)

- **Shipped:** Omnichannel ticketing (Slack, Teams, WhatsApp, Discord, email, chat, forms); AI Agents (Runbook-based deflection, ticket triage, intelligent routing, contextual draft replies, autonomous field/status updates); AI Assistants (in-Pylon productivity, "3x speed"); Account Intelligence (account health scoring, churn risk detection, multi-stakeholder views); AI-powered KB (auto-drafted articles, gap detection); AI Product Intelligence (auto-cluster feature requests with ARR weighting, Linear/Jira links); CRM integrations (Salesforce, HubSpot, Linear, Jira); SOC 2 Type II, ISO 27001, HIPAA, GDPR.
- **Goals orientation:** None. Account Intelligence is the closest analog but is metric-based, not goal-declarative.

- **QA rigor:** Light vs Decagon, Sierra, Intercom; no native AI scoring or QA product disclosed.
- **B2B specific workflows:** Slack Connect as a first-class channel; shared customer workspaces as structured tickets; multi-stakeholder account views; SLA management; per-account context.
- **Pricing:** Per-seat subscription; no public pricing; sales-led demo only.
- **Momentum:** \$31M Series B (January 2025) co-led by a16z + Bain Capital Ventures; \$51M total raised; 750+ customers; 150+ migrations from Zendesk, Intercom, Salesforce Service Cloud; 5x+ YoY revenue growth for 2 years; logos include Deel, Cognition, Together AI, Temporal, AssemblyAI, Ada, Hightouch.
- **B2C vs B2B:** Pure B2B.
- **Threat level: High specifically for Kustomer's B2B expansion.** Pylon is the de facto standard for AI-native B2B post-sales tooling.
- **Kustomer wedge:** Bring goals-first config and AIC 2.0 agent rigor to B2B accounts. Pylon has not connected account intelligence to a goals primitive.

Comparison summary

Vendor	Goals-first orientation	Threat & wedge for Kustomer
Intercom Fin	Vision only. Pioneer 2025 promised "Goals"; not shipped. Closest analog is Fin for Sales "Playbook" (role-scoped).	Highest. Most likely to ship a goals primitive in next 6 to 12 months. Wedge: Pro paywall (\$99/mo) gates analytics; Kustomer bundles into Beacon.
Zendesk	Outcome-based pricing, no admin-declared goals. Custom QA is post-hoc, not goal-led.	High from scale + Forethought (closed March 2026). Wedge: Zendesk is built around channels; Kustomer can be built around goals.
Sierra	Closest philosophically. Reality: per-agent guardrails, not workspace goals.	High at enterprise; \$150K ACV floor limits mid-market. Wedge: Kustomer covers AI + human + reporting in one model.
Decagon	Per-AOP outcomes; Watchtower QA and Experiments are best-in-class. No workspace-level goals.	Medium-high. Strong engineering depth, weaker brand. Wedge: Decagon is agent-engineering-led; Kustomer is goal-led for business users.

Vendor	Goals-first orientation	Threat & wedge for Kustomer
DigitalGenius	None at product level.	Low-medium. Narrow to fashion/DTC. Wedge: Kustomer is horizontal with goals as the unifying primitive.
Pylon	None. Strong account intelligence, but goals not declarative.	High for B2B. 750+ customers, 150+ migrations from Zendesk/Intercom/Salesforce. Wedge: bring goals-first config to B2B; Pylon has not connected health scoring to goals.

Appendix B: QA Platform Competitive Summary

The QA platform category is the most likely place Beacon will be compared to. Important context: Beacon is **not** a QA platform. Beacon is a goals-first governance product where Quality Monitors are one of several components. This appendix exists so we can frame the conversation correctly when prospects compare us to standalone QA tools.

MaestroQA

- AI Platform on Amazon Bedrock (2024 to 2025) for advanced conversation analytics. Integrations with Sierra, Decagon, Agentforce, Forethought (positioning as the QA layer atop any AI agent).
- No goals primitive.
- Per-evaluation/per-seat pricing; sales-led.
- **Threat to Kustomer:** Low directly, but high as a signal that QA-on-AI-agents is becoming the new standard. Beacon must clearly differentiate as goals-led, not QA-led.

Level AI

- iCSAT (January 22, 2025) auto-scores CSAT on 100% of interactions. Unified Intelligence Loop (February 2026). AgentGPT (2023), Auto-QA, Voice of Customer, coaching dashboards.
- No goals primitive; correlates KPIs with iCSAT but does not let admins declare workspace goals.

- Enterprise contact-center focus.
- **Threat to Kustomer:** Low-medium. Strong in voice; not a horizontal CX platform.

Observe.AI

- VoiceAI Agents (GA March 26, 2025); AutoQA on 100% conversations; Agent Copilot, Coaching Copilot, Insights Copilot; 25+ languages; Real-Time AI; integrates with Salesforce, Zendesk, ServiceNow; HIPAA, SOC 2, ISO 27001. Acquired Dubdub.ai (March 2025).
- No goals primitive.
- **Threat to Kustomer:** Medium for voice contact centers, less for digital CX.

Scorebuddy

- GenAI Auto Scoring (August 2024); 100% interaction coverage; fully customizable scorecards. AI Knowledge (RAG-grounded scoring against KB). **Intercom partnership** (February 2025); Intercom itself uses Scorebuddy for internal QA.
- No goals primitive.
- **Threat to Kustomer:** Low directly. Telling that Intercom uses Scorebuddy rather than build internally — confirms QA is genuinely hard and that AI vendors are choosing to partner rather than build it themselves.

NICE CXone (Playvox)

- **NICE acquired Playvox October 1, 2024** for approximately \$40M. QM Auto Score GA February 2026. Consolidates QA/WFM/EE.
- No goals primitive; NICE Enlighten AI provides outcome correlations.
- Integration risk for existing Playvox mid-market customers (Calabrio analyst commentary flags potential migration and pricing pressure).
- **Threat to Kustomer:** Low for Kustomer's segment. NICE owns high-end voice contact center; minimal overlap with Kustomer's digital CRM.

Why Beacon is not a QA platform

QA platforms grade conversations after they happen. Beacon grades **goal progress**, of which conversation quality is one input. The distinction matters because QA platforms cannot influence AI agent configuration, rep assignment, or admin setup — they observe. Beacon influences all of those because Goals are a first-class object in the Kustomer platform. This is the narrative Sales and Marketing must consistently land.

Appendix C: Adjacent Inspirations

These are not competitors. They are products from adjacent categories whose primitives and UX patterns Kustomer should study and borrow from. One per category, picked for the brand recognition that lets Kustomer execs and the broader org relate quickly.

Observability / SRE: Datadog SLOs

Datadog's Service Level Objective product is the canonical "declare a target, the platform tracks, alerts, and reports against it" pattern.

- **Primitives:** SLI (the metric), SLO (target percentage over window), error budget (how much breach is acceptable), burn rate (rate of error budget consumption), alert.
- **Standard windows:** 7d, 30d, 90d.
- **What Kustomer should borrow:**
 - **Error budgets,** not binary pass/fail. A goal of "CSAT $\geq$ 4.5" is more useful framed as "we have an error budget of N conversations below 4.5 over 30 days."
 - **Burn-rate alerts:** proactive when goal trajectory is at risk, not just when goal is missed.
 - **SLO widgets on dashboards:** any stakeholder can see goal status inline with operational data.
 - **SLO-as-Code (Terraform pattern):** enterprise admins can version goals in source control.

Revenue Intelligence: Gong

Gong's product philosophy: capture nearly every customer interaction, run AI scorecards (AI Call Reviewer released January 27, 2026), tie conversation behavior to deal outcomes via Gong Deals + Gong Forecast, and surface "AI Tracker" patterns that correlate behavior with revenue.

- **What Kustomer should borrow:**
 - **Outcome-back-correlated scoring:** which agent behaviors actually correlate with goal achievement, not just rubric adherence. (Most QA platforms score against rubrics without ever checking which rubric items correlate with outcomes.)
 - **Coaching workflows tied to evaluation:** Gong's scorecards drive structured coaching loops. AIR should do the same.
 - **One scoring primitive across AI and humans:** Gong scores reps and AI assistants on the same scale. Beacon's Quality Monitors should do the same across AIC and AIR.

Product Analytics: Amplitude North Star Framework

Amplitude's North Star Metric framework: one metric reflecting customer value, leading indicator of revenue, broken into 3 to 5 input metrics. Their own NSM evolved from Weekly Querying Users to Weekly Learning Users with three inputs (Activated Accounts, Broadcasted Learnings, Consumption of Learnings).

- **What Kustomer should borrow:**
 - **Goals as a tree:** lagging outcome decomposes into 3 to 5 leading input metrics. A goal like "improve CSAT by 5 points" decomposes into specific drivers.
 - **North Star workshop UX as goals-first onboarding metaphor:** Amplitude published their workshop playbook. It is the right model for Beacon's goals-first admin setup.
 - **Outcomes ≠ outputs:** refuse to let admins set goals on volume or activity. Only on real customer outcomes.

AI Evaluation: Braintrust

Braintrust's evaluation primitives: scorers (code, LLM-as-judge, human), datasets built from production traces, experiments that compare prompt/model versions, CI/CD gates that block regressions before deployment, and "Loop" (an AI assistant that helps non-technical users analyze logs and generate evals in plain language).

- **What Kustomer should borrow:**
 - **Production traces auto-convert into eval cases:** every escalated AIC 2.0 conversation should be reusable as a regression test.
 - **LLM-as-judge scorers with explicit rubrics:** the AIC 2.0 Validation Agent should be exposed as a first-class scorer type.
 - **Scorer transparency:** "why this conversation was flagged" must be exposed for admin trust.
 - **Evaluation blocks bad releases:** Beacon's Quality Monitors should be able to block an AIC configuration change if it regresses on a goal.

Appendix D: Risks and Open Questions

Risks

1. **Intercom ships Goals first.** Their CEO has committed publicly. If they ship in the next 6 months, our first-mover claim weakens.
Mitigation: accelerate Beacon GA, ensure analyst briefings position goals-first with Kustomer before Intercom's launch.

2. **Parallel-tracking fractures the primitive.** If AIC, AIR, and AI for Leaders teams each invent their own "goal" model, the vision is incoherent. Mitigation: Beacon team owns the Goals schema; all other teams consume it.
3. **Admin overwhelm.** Asking admins to declare goals during setup is a real friction. Mitigation: ship strong defaults (5 to 7 OOTB goals); admins can accept, edit, or skip.
4. **Goal misuse.** Admins set vanity goals (volume, activity) that do not drive business outcomes. Mitigation: enforce a small taxonomy of goal types tied to real customer outcomes. Reject activity metrics as goals.
5. **B2B vs B2C goal divergence.** B2C goals (CSAT, deflection) and B2B goals (NRR, renewal saves) have different data models. Mitigation: Goals schema accommodates both from day one (already scoped in Beacon M1).

Open questions

- How many goals per org? Cap or unlimited?
- Goals scoped to specific AI agents or always workspace-level?
- Goals tied to specific accounts/segments, or only across all conversations?
- Does Kustomer want goals to be addressable from outside Kustomer (Snowflake, Looker)?
- Is there a separate pricing motion for goals-first capabilities, or are they bundled into the core platform?

End of doc.