

ANALYSIS PREPARED BY YVONNE DOLL, MAY 21 2026

A goals-first IA for the whole Kustomer platform.

Goals and Performance become core platform destinations alongside Inbox. Lower-frequency items move to overflow navigation, while the AI workspace focuses on one job: simplifying AIC and AIR configuration around organizational goals.

01 Principles for the goals-first IA

- 1 **Goals are platform primitives.** They live in the site rail. Every other destination references them.
- 2 **Outcome on Goals, operational on AI.** Each goal shows its trajectory. AI → Analyze shows whether the automation is healthy. Performance is the cross-goal operational rollup.
- 3 **Setup, not configuration.** AIC and AIR each get one opinionated setup surface. Defaults loud, advanced knobs quiet.
- 4 **One wizard, two branches.** Get Started asks AIC, AIR, or both, then runs the right sequence end to end.
- 5 **Goal context follows the user.** Conversations, alerts, suggestions, and report rows know which goals they touch and link back.
- 6 **Automations are first-class and shared.** Procedures, tools, knowledge, and shared guardrails in Building Blocks are reused across automations. Edits propagate.
- 7 **Per-automation settings live with the automation.** Competitors, Secrets, Tone, greetings, verification, routing live on each automation's Settings. Shared guardrails (refund ceilings, escalation triggers) live in Building Blocks and attach inline.

- 8 **Topics are org-level.** One canonical taxonomy under Reporting → Topics, auto-detected from history. Goals select a subset; everything else references the same list.

- 9 **Goals facet by audience, not by surface.** Every goal carries an audience tag (AIC, AIR, Sales, Human Reps, Other). The Goals destination filters by audience; AI Setup pre-filters to AI goals.

- 10 **Advanced is preserved, not promoted.** Build, Test, Deploy, Manage Automations get a LEGACY pill under a "For power users" subheader.

- 11 **Recommendations apply in place.** Every suggestion carries an "Apply in [block]" CTA that opens the affected procedure / KB / tool / guardrail inline. No page jumping.

- 12 **Inline capture, canonical home, persistent provenance.** Capture new data where it's first needed (Company profile, Topics, audience tags). Store at its canonical org-level home. Stamp the home on every consuming surface with an edit link.

- 13 **One lifecycle vocabulary: Build / Test / Deploy / Analyze.** Same verbs in AI Setup, per-automation tabs, suggestion CTAs.

- 14 **One Build tab. First-run is a mode of it, not a separate destination.** Creating a new automation routes the admin to the same per-automation Build tab they'd edit on day 90, with a stepper, locked future tabs, and parser-drafted callouts overlaid until Deploy activates. No parallel "wizard" surface.

- 15 **Build and Test share a right rail: Assistant + Test Console.** Persistent, collapsible. Assistant chats about the open guidance; Test Console sends live messages on a chosen channel and brand. Same rail, same widget contract on both tabs.

- 16 **Procedures have associations, not owners.** Every procedure declares which surfaces it applies to: scope labels ("All AI for Customers", "All AI for Reps") and specific automation ids. The library is the canonical home; each automation's Build shows its slice. Editing a multi-scope procedure fires a blast-radius prompt.

- 17 **Deploy is cross-audience by design.** One page with audience banners: an AIC section (conditions, smart routing, eval recap) and an AIR section (Copilot toggle, mode, teams). Same Activate button covers both. AIR Settings still owns steady-state Copilot config; Deploy owns the activation moment.

- 18 **The recommendation→action loop is structural, not a link.** Every suggestion names its target block. The loop has its own affordance, its own pattern.

- 19 **Three metric layers, one question each.** Goals = "did we hit the outcome?" · Performance = "are agents healthy right now?" · Analyze = "is this automation behaving?" Each surface declares its layer in chrome; drill direction is one-way.

02 Explore the proposed IA

Click any node to read its purpose, the rationale for placing it here, and a sketch of what its screen contains. The tree starts at the Kustomer platform level so the introduction of Goals and the elevation of Performance are visible at a glance.

KUSTOMER PLATFORM, IA TREE

Expand all

Collapse

🏠

Kustomer Platform

▼

🎯

Goals

NEW ▼

• Create Goals (slide-out panel)

NEW

• Goal detail

NEW

📈

Performance

NEW >

🌟

AI

EXISTING >

🕒

Reporting

EXISTING >

🏠

Rep Dashboard

RENAMED

📧

Inbox

EXISTING

🔍

Searches

EXISTING

📈

TeamPulse

EXISTING

🔍

Global Search

EXISTING

Kustomer Platform > Goals

Goals

Cross-goal outcome rollup at the top, then the list. "Are we hitting what we said we'd hit?" Sparkline grid of every goal, at-risk lane, contribution-by-audience aggregate. Leadership view.

Goal 1

LIVE

Reduce refund tickets 20%

**Goal 2**

TESTING

Speed up rep responses

**Goal 3**

ATTENTION

Save at-risk renewals 80%

**PURPOSE**

The org's declared outcomes plus the leadership-level read on whether they're being hit. The header is a cross-goal outcome summary; the body is the audience-faceted goal list. Distinct from platform Performance, which is the cross-agent operational rollup.

 GOALS REFERENCE

This is where Goals are declared and tuned. Every other destination reads from this primitive, pre-filtered by audience when context is implied. The slide-out form is the single canonical authoring surface for a goal regardless of where it opens from.

 TOPICS REFERENCE

Goals select a topic subset to declare scope. Topic chips on each goal define what conversations it covers; coverage rollups read from the canonical Topics taxonomy.

 SUGGESTIONS REFERENCE

AI Suggestions surface as ranked recommendations to fix what monitors and anomalies surface against each goal. Cross-goal queue lives at Performance, Suggestions.

 OTHER TOUCHPOINTS

- **AI Setup**, provides scope (audience, topics, scenarios) for every new automation
- **AI for Customers / AI for Reps**, every attached automation drives one or more goals

- **Performance**, outcome trajectories roll up across all goals
- **Reporting → Data Explorer**, goals are a first-class dimension in ad-hoc queries
- **Performance, Suggestions**, suggestions targeting goals surface here

WHAT'S IN HERE

- **Outcome summary header** at the top of Goals Home: cross-goal sparkline grid, at-risk lane, contribution-by-audience aggregate (AIC / AIR / Human / Sales / Other), % on track / at risk / draft
- Links from any summary tile into platform Performance scoped to the contributing AI / human signals
- Goal list view with state, audience, owner, and progress
- Audience facet filtering and role-default scope
- State filters (Live, Limited, Testing, Attention, Draft)
- Time-window selection for trajectory readings
- Goal templates entry point
- Goal reports, per-goal trajectory, contribution breakdown by audience (AIC / AIR / Human / Sales / Other), period-over-period delta, topic-level slice, export, and schedule
- Cross-team goal authoring and management
- Audience faceting (AIC, AIR, Sales, Human Reps, Other)
- Goal templates and starter library
- Per-goal detail pages with outcome trajectory
- Goal pre-filtering by viewer role
- Pre-filtering contract for goal pickers opened from any destination

WHAT YOU CAN DO

+ Create goal

opens Create Goals slide-out

Browse templates

6 OOTB goal templates

Filter by audience

AIC · AIR · Human Reps · Sales · Other

Filter by state

Live · Limited · Testing · Attention · Draft

Open goal detail

click any row to drill in

Time-window

7d / 30d / 90d trajectory readings

Open report

per-goal contribution breakdown

[See cross-goal Performance](#)

→ Performance scoped to contributing signals

WHY IT LIVES HERE

Lifted out of the AI workspace and reframed as a platform primitive. The faceting model is the unlock that lets every team's goals coexist in one list without forcing AI admins to wade through sales-pipeline targets, or sales leaders through deflection percentages. Authoring stays a single canonical primitive (one slide-out); browsing adapts to context via persistent audience filters.

03 The full CX, end to end

One concrete journey for an admin declaring an org-level goal, attaching the AI, human workflows, and monitors that drive it, and tuning the loop a month later. Goals are the spine, AI is one of several things attached to them. Every step lands on a destination that exists in the IA above; nothing is invented mid-flow.

0 First-run primer ONE-TIME · ~30 SECONDS · DISMISSABLE

The very first time an admin opens the platform, a stepped mini-flow introduces the new primitives. Each sub-step is a short card with what it is, where it lives, and how to get there. Skippable.

0.1**Goals are the spine**

Every AI automation, monitor, and human workflow connects to a goal. Goals organize the platform around outcomes.

Lives in **Goals** (site rail)

**0.2**

Topics, auto-detected

The shared taxonomy of what customers ask about. Auto-detected from the last 90 days. Editable.

View: **Reporting** → **Topics**
Manage: **Settings** → **Topics**



0.3

Suggestions apply in place

When AI proposes a fix, it surfaces where the affected block lives. Apply, edit, or dismiss inline.

On every block + **Performance** → **Suggestions**



0.4

Company profile, captured inline

During your first goal, we'll ask about industry, channels, brands. Stored once, read everywhere.

Captured at **Goals** → **Create goal**
Canonical home: **Settings** → **Company profile**



For existing AI customers, a fifth migration sub-step (0.5)

Lists per-automation goals found in the workspace with a one-click **Lift these into platform Goals** action. Topics taxonomy is seeded automatically from conversation history; admin reviews and approves.

1 **Go to Goals** PRIMARY NAV

Click Goals in the site rail. Goals Home opens pre-filtered to the viewer's role (AIC for this admin), with the 6 OOTB templates as a lead-in.



2 **Create your first goal & capture Company profile inline** GOALS → CREATE GOALS SLIDE-OUT

Click **+ Create goal**. The slide-out opens with a **Company profile** section at the top (industry, channels, brands, description) — captured inline on first run because it's needed for smart goal suggestions. The same form captures the goal itself: *"Reduce refund tickets by 20%",* audience *AIC,* owner self, target *20% over 90 days,* topics (multi-select from Reporting → Topics), and the org's **rep-training doc / SOP**. Save commits the Company profile to **Settings → Company profile** (canonical home, edit there anytime) and creates the goal. Subsequent goals collapse the Company profile to a one-line context strip with an "Edit in Settings" link.

3 **Land on goal detail** **GOAL DETAIL**

Topic chips and coverage stats appear at the top ("3 of 5 topics, 54% of 90d volume"). The parser surfaces draft procedures + suggested tools + knowledge from the uploaded doc, ready to attach. Empty attached-resources state below invites attaching an AI automation, monitor, or human workflow.

4 **Attach AI automation** **GOAL DETAIL → + ATTACH AI AUTOMATION**

Click **+ Attach AI automation**. Name it *Refund Order*. A new automation row is created in **Automations** with `firstRunMode: true`, and the admin lands on the new automation's **Build** tab. Same surface as steady-state editing, in first-time mode (stepper + locked future tabs + parser callouts).

5 **Build, in first-time mode** **AUTOMATIONS → REFUND ORDER → BUILD**

Stepper at the top: ● **Build** → ○ Test → ○ Deploy → ○ Analyze (future tabs locked). Goal-context strip at the top (topics, scenarios). Capture only AI-specific knobs: *Where the AI shows up* (pre-set to AIC) → *Knowledge sources* → *Tone of voice*. Parser-drafted procedures, knowledge, and tools surface with "AI added" callouts, inline-editable. **Save & continue** unlocks Test.

6 **Test** **AUTOMATIONS → REFUND ORDER → TEST**

Stepper: ● Build → ● **Test** → ○ Deploy → ○ Analyze. Click **Run All Tests**. The AI runs against sample conversations grouped by category, Standard Refunds, Returns, Late Delivery, etc. Same Test tab the admin would use on day 90, with the stepper overlay.

7 **Analyze test results** **AUTOMATIONS → REFUND ORDER → TEST (RESULTS VIEW)**

5 pass, 1 fails, *Damaged or Defective Items* at 23%. Three ways forward.

THREE WAYS FORWARD

Fix it now

Jump back to **Build** via the stepper. Refine procedures, knowledge, or tone inline. Re-test. Same surface, same automation row.

Deploy partial

Activate, but exclude the failing category. Smart-routing rule auto-generated to hand it off.

Override

Proceed with confirmation. Failing category flagged on the goal page so it isn't forgotten.

→ CONTINUE TO DEPLOY

8

Deploy & Activate

AUTOMATIONS → REFUND ORDER → DEPLOY

Stepper: ● Build → ● Test → ● **Deploy** → ○ Analyze. Set conditions (channel, body filters), confirm smart routing, click **Activate**. On activate, **first-time mode ends**: the stepper vanishes, all tabs unlock, and the automation behaves like any other.

9

Back to goal detail, AI live

GOAL DETAIL

Goal shows **Refund Order: Active**. Performance section transitions from empty to live; trajectory begins. Outcome performance lives here on the goal; operational signals on the AI itself live in AI → Analyze.

10

Two weeks later, a suggestion appears

SUGGESTIONS & ALERTS

Goal shows **-14%**, on track. The Validation Agent has drafted **Suggest an upsell before applying a coupon code: High confidence (82%)**, tagged *Increase CSAT*. The card shows **why this may matter**, **how the pattern was detected** (33% of coupon requests came from customers who hadn't reviewed related products), the supporting **2 conversations failing the same way**, and the **before / after** diff inline. Three ways to act.

THREE WAYS TO ACT ON A SUGGESTION

Apply

One click commits the AI-drafted diff. Slide-out closes; change is live on the procedure / KB / tool / guardrail. Version history records "applied from suggestion".

Edit & apply

Click **Edit Procedure**. Slide-out opens with the affected block pre-loaded and AI changes marked "AI added". Tweak inline, then **Save and apply**. No page navigation.

Dismiss

Reject with a reason. Feeds back into the model so similar suggestions don't recur. Audit trail captured.

→ CHANGE IS LIVE, PERFORMANCE KEEPS TRACKING

11

A month later, extend

GOAL DETAIL, AUTOMATIONS

From the goal page attach a **performance monitor** (Quality Monitor scorecard) and a **human workflow** for edge cases. If the goal needs a second AI flow ("Damaged Items"), click **+ New automation** on Automations or **+ Attach AI automation** from the goal, then go through Build / Test / Deploy in first-time mode again.

04 The recommendation→action loop

The structural pattern that closes the gap between detection and fix. Every AI suggestion carries an explicit *Apply in [block]* affordance that opens the exact procedure, monitor, KB

article, tool, guardrail, or coaching path being recommended, pre-scoped and ready to edit. The loop has its own affordance, its own pattern, its own home in the IA. It is the differentiator, not a footer link.

STAGE 1, SIGNAL

Something drifted

Monitor scorecard fails or trends down
Validation Agent flags an anomaly
Goal trajectory crosses an at-risk threshold
Pattern detected across a cluster of conversations

Refund eligibility check

AIC v3 · Pass rate 62% · ↓14% over 7 days

STAGE 2, SUGGESTION

A draft fix appears, where it belongs

Card surfaces on the affected block + in the cross-AI Suggestions queue
Confidence score, goal tag, before / after diff inline
Why this matters, how the pattern was detected, supporting conversations
The card explicitly names its target block in the CTA

Suggest an upsell before applying a coupon

High confidence (82%) · drives *Increase CSAT*

→ Apply in Procedure: [Coupon Handling](#)

STAGE 3, ACTION

The block opens in place

Inline slide-out opens with the block pre-loaded
AI-proposed changes marked "AI added"
Tweak, apply, or dismiss with reason, never leaving the page
Version history records "applied from suggestion #N"

Procedure: Coupon Handling

+ Suggest upsell before code (AI added)
– Existing: validate code, then apply

[Save and apply]

THE SIX TARGET BLOCKS A SUGGESTION CAN OPEN

→ Apply in Procedure

Opens the procedure body with AI changes inline; lives in Building Blocks, attached to the automation.

→ Apply in Monitor

Opens the Quality Monitor scorecard criteria with weight or threshold tweaks pre-staged.

→ Apply in KB Article

Opens the Knowledge Source article with AI-drafted content additions or corrections.

→ Apply in Tool

Opens the tool config (auth, scope, parameters) with the proposed change pre-set.

→ Apply in Guardrail

Opens the per-automation guardrail or the shared guardrail rule with the new constraint pre-filled.

→ Apply in Coaching Path

For human-rep degradation: opens the coaching path with the targeted skill or scenario pre-selected.

05 Risks the new Goals primitive introduces

The strategic costs of treating Goals and Topics as platform primitives and splitting outcome performance from operational performance. Severity reflects blast radius across personas.

#	RISK	SEVERITY	WHY IT HURTS
1	Rail real estate	High	Goals and Performance newly promoted alongside the existing rail (Inbox, Searches, Reporting, TeamPulse, Rep Dashboard). Rail runs eight items deep before utility; Apps, KB, Widgets, Settings, Help all get pushed into overflow.

#	RISK	SEVERITY	WHY IT HURTS
2	Cross-team primitive ownership	High	One Goal object. One edit surface. One set of display primitives with contextual variants — every surface picks a variant, no surface draws its own.
3	Audience + topic tag drift	Medium	Mis-tagged goals or conversations shatter every picker and rollup. Required-on-create, auto-detect confidence flags, and audit jobs mitigate.
4	Audience boundary ambiguity	Medium	"Improve CSAT" belongs to AIC, AIR, and Human Reps at once. Multi-select audience + primary-owner field + documented rubric.
5	Goal reports left Reporting	Medium	Goal-specific reports moved to Goals Home. Reporting users may not find them. Cross-link callouts from Data Explorer and Standard Reports.
6	Goal-attribution noise	Medium	Every block, toggle, monitor, and suggestion carries "drives Goal X" annotations. Without density rules they become wallpaper. Show only when scope is implied; cap per surface.
7	Primer skip + migration mis-mapping	Low	Admins skim primers. Mitigation: Company profile is captured inline during first goal creation (can't be skipped); make the primer re-openable from the avatar menu; the auto-lift of per-automation goals into platform Goals runs a one-time audit on migrated goals.
8	"Override" exit on Analyze	Low	The third branch (Override) ships an automation with a known-failing category. Goal page must flag it loudly so it isn't forgotten.
9	"Resources" renamed to Building Blocks	Low	Existing admins navigate by muscle memory. Add an inline "formerly Resources" hint for one release; redirect deep links.
10	Per-automation surface scope creep	Low	The per-automation Settings page could grow back into a multi-tab spread. New

#	RISK	SEVERITY	WHY IT HURTS
			toggles default to advanced until usage promotes them.

06 Friction points and resolutions

Each remaining tension in the proposed IA, paired with the standing contract that resolves it. Every contract below is documented in the relevant IA node so the resolution rides with the structure.

FRICTION	WHAT'S HAPPENING	CONTRACT (DOCUMENTED IN IA)
Goal authoring entry points diverge	" + Create goal" surfaces in six places: Goals Home, the Attached Goals card on Build, per-automation Settings, Manage Automations, AI Setup, and every goal-picker footer.	One authoring component, the Create Goals slide-out. Every entry point opens it with audience + topics pre-filled and returns the saved goal pre-selected.
Outcome rollup vs operational rollup look alike	Goals Home opens with an outcome summary ("hitting what we said?"). Performance is the operational rollup ("how are agents doing?"). Both look like dashboards.	Distinct framing copy in each header; never duplicate tiles. Goals Home tiles deep-link to Performance; Performance rows link back to the goal.
Same primitives at platform and per-automation	Monitors, Suggestions, Anomalies live as Performance children and as Analyze tabs inside every AIC automation. Same primitives, two scopes, unclear which to act on.	Per-automation tabs own authoring (apply, dismiss). Performance is the aggregate. Rollup rows carry a scoping pill that deep-links to the source; sources footer-link to "across all automations".
Anomalies is AI-only inside a cross-agent rollup	Performance's three children look like peers, but Anomalies (Validation Agent) is AI-only; only Monitors and Suggestions cover human work.	"AI only" subhead on Anomalies + one-line explainer on Performance. Human degradation surfaces via Monitors (dips) and Suggestions (coaching).
Guardrails split into two flavors, one word	Per-automation guardrails (Competitors, Secrets, Tone) live on each automation's Settings. Shared guardrails (refund ceilings,	Rename in chrome: "Per-automation guardrails" on Settings, "Shared guardrails" in Building Blocks. Inline

FRICTION	WHAT'S HAPPENING	CONTRACT (DOCUMENTED IN IA)
	escalation triggers) live in Building Blocks. Same label, two homes.	" + Attach shared guardrail" on each automation's Settings.
"Mine" vs audience facets on Goals	"Mine" overlaps with audience facets, an AI admin's "Mine" is almost always their AI goals.	Pre-select the audience facet by viewer role; "Mine" filters within whatever facet is active.
Topics management vs Topics analytics	Authoring (rename, merge, hide, custom instructions) lives in More → Settings → Topics. Analytics lives in Reporting → Topics. Same name, two homes.	Keep each page single-purpose; both already cross-link. Never duplicate controls across them.
Computed Fields moved out of Building Blocks	Computed Fields now live in More → Settings. AI admins who used to find them in the AI workspace may not discover the new home, and every goal threshold reads one.	Cross-link every goal threshold and scorecard into Settings → Computed Fields, field pre-selected. Reverse-link each field to the goals and monitors it feeds.
AI Setup wizard duplicated per-automation tabs	Old IA had AI Setup as a parallel surface authoring the same data the per-automation Build / Test / Deploy / Analyze tabs already capture, with the same verbs since the lifecycle-vocabulary rollout.	Resolved. AI Setup is removed. " + New automation" creates a row in Automations and routes to its Build tab in first-time mode (stepper, locked future tabs, parser callouts). Same URL, same surface, one place to author and edit.
AIC selector vs Manage Automations	Clicking AI for Customers opens a switcher dropdown; Manage Automations is a separate destination. Both show the AIC automation list.	Dropdown is switch-only: name + status pill + " → Manage all" footer. Manage Automations owns add, edit, delete, duplicate, filter, sort, bulk.
Tools not visible in primary nav	Shared Tools are the AI's hands, nesting them inside Building Blocks understates their operational weight when an integration breaks.	Flyout on the Building Blocks rail row exposes Procedures, Knowledge Sources, Tools as one-click destinations.

07 Jobs to be done (Usability test)

Fifteen testable scenarios across the three pillars of the goals-first thesis. Each task is framed as a participant-facing scenario with a success signal a moderator can observe. Sources: Goals-Based AI Monitoring (Aditya, Apr 2026), Goals-Based AI Setup (Amisha, May 2026), and this IA recommendation.

Goals

Goals as a platform primitive. Authoring, scoping, leadership rollup, and goal-context-follows-the-user.

Monitors

Quality + Anomaly monitoring of live AI. Per-criterion drill-down, alerting, suggestions, runtime QA.

AI Config Setup

Goals-first intake, plan trust, staged testing, scoped deploy, closed-loop recommendations.

Goals Maps to monitoring JTBD 1, 4, 8 + the leadership rollup pattern

5 TASKS

1 Declare a new business goal from scratch

"You've been asked to reduce refund-related escalations. Set up a goal to track this."

SUCCESS Participant lands on Goals, chooses Create, sets a target value, direction (lower is better), and unit without backtracking.

TESTS Goals as primitive, slide-out authoring, defaults vs. custom.

2 Scope a goal to a specific AI agent

"Your refund agent is one of three AI agents. You only want this goal to track that agent's performance."

SUCCESS Participant finds the Goal Resource control and selects the correct AIC automation.

TESTS Resource scoping, goal-to-automation linkage, whether the relationship is discoverable.

3 Read the leadership rollup, identify an at-risk goal

"You're prepping for a Monday exec review. Tell me which goals are on track and which need attention."

SUCCESS Participant uses Goals Home (cross-goal sparkline grid) to identify the

4 Trace a goal back to the conversations driving its score

"Your CSAT goal is dipping. Show me the conversations contributing to last week's drop."

SUCCESS Participant drills from the goal into the linked monitor's flagged

- TESTS

at-risk goal within 30 seconds and can articulate why.

Outcome rollup legibility, at-risk lane, sparkline comprehension.
- TESTS

conversations.

Goal → Monitor → Conversation traceability, the "goal context follows the user" principle.

5 Edit a goal's target mid-flight

"Leadership has raised the bar, the team now needs to hit 90% CSAT instead of 85%. Update the goal."

- SUCCESS

Participant finds the goal, opens it, edits the target, and saves without confusion about whether monitors or alerts need to be reconfigured.
- TESTS

Edit affordance, whether dependencies are surfaced, optimistic vs. anxious editing behavior.

Monitors Maps to monitoring JTBD 2, 3, 5, 6, 7, 9

5 TASKS

1 Find out why a specific conversation was flagged

"This conversation was flagged by your Quality Monitor. Tell me which criterion it failed and why."

- SUCCESS

Participant opens the conversation, sees per-criterion scores, weights, pass/fail, and rationale.
- TESTS

Per-criterion drill-down (JTBD 5), readability of the rubric in context.

2 Configure a custom monitor with weighted criteria

"You want to start scoring AI conversations on tone and accuracy. Tone is twice as important as accuracy. If accuracy fails, the whole conversation fails."

- SUCCESS

Participant creates a custom monitor, adds two criteria, sets weights, marks accuracy as essential, sets a pass threshold.
- TESTS

Criteria authoring (JTBD 9), the essential flag concept, weight comprehension.

3 Opt into alerting for a degrading monitor

"You want to be told within the hour if this

4 Act on an AI-generated suggestion from a failed monitor

"The monitor has produced a suggestion to update this procedure. Review it and either

monitor's pass rate drops. Set that up."

- SUCCESS** Participant turns on alerts, selects hourly cadence, confirms delivery channel.
- TESTS** Opt-in alerting (JTBD 3), cadence selection, delivery channel clarity.

apply or dismiss it."

- SUCCESS** Participant finds the suggestion card on the Suggestions & Alerts surface, reads confidence + goal tag + the **before / after** diff inline, clicks **Edit Procedure** to open the inline slide-out with the affected procedure and AI-marked changes, applies (or edits then saves) or dismisses with a reason. No page navigation.
- TESTS** Suggestion legibility (JTBD 6), before/after pattern, the inline slide-out edit-in-place affordance, dismiss-with-reason.

5 **Review anomalies surfaced by the Validation Agent**

"Your AI agent went off-script three times last week. Find those events and tell me what happened."

- SUCCESS** Participant locates the Anomaly Detection Report, identifies the three events, and reads the anomaly type and affected resource.
- TESTS** Anomaly Detection discoverability (JTBD 7), the connection between anomalies and suggestions, runtime QA as a separate concept from post-call QA.

AI Config Setup

Maps to the four JTBDs in the Goals-Based AI Setup doc + the 5-page flow

5
TASKS

1 **Set up Kustomer AI for the first time, anchored on an outcome**

"You just signed with Kustomer. You're responsible for reducing escalations on refund conversations. Get your AI configured."

- SUCCESS** Participant completes the intake (company context, audience, goals,

2 **Review and trust the generated plan before saving**

"Kustomer just generated a plan based on your answers. Walk me through what it's going to do, and tell me whether you'd save it as-is or change something."

- SUCCESS** Participant can name what was generated (procedures, monitors,

TESTS

topics, monitors, knowledge, tone) without abandoning. Articulates that they're configuring around the escalation goal, not assembling parts.

TESTS

The core thesis: goals-first intake feels different from configuration-first. Measure abandonment, backtracking, and post-task sentiment.

TESTS

alerts, scope), identify which goal each item ties to, and edit at least one item.

TESTS

Plan traceability (JTBD 2, "see what will change, where, and why"), blast-radius comprehension, editability of generated content.

3 Test the staged plan before deploying

4 Deploy with scope and capture a deploy note

"Before this goes live to customers, you want to make sure the AI handles refunds correctly. Test it."

SUCCESS

Participant finds the Test surface, runs an evaluation category against the staged plan, and interprets the readiness signal. Optionally uses the side panel test console for an ad-hoc conversation.

TESTS

Discoverability of staged-vs-production distinction, readiness signal comprehension, the separation of save from activate.

"You're ready to turn this on, but only for chat conversations from your US brand. Deploy it."

SUCCESS

Participant scopes via rule groups (channel + brand), writes a deploy note, confirms.

TESTS

Scope-then-activate pattern, rule group comprehension, whether the note feels valuable or like friction.

5 Act on a post-deploy recommendation, inline

"Two weeks in, your CSAT goal is drifting. Kustomer is recommending a procedure tweak. Apply it."

SUCCESS

Participant finds the suggestion where it lives in context (on the procedure card / in the suggestions queue), clicks **View suggestion**, the recommendation slides out with an inline edit view of the affected procedure, they apply or edit and confirm without leaving the current page.

TESTS

The closed-loop story (JTBD 3, "tell me which lever to pull"), inline slide-out edit pattern (no page jump),

preview-before-apply, version
history awareness.

Prepared for Amisha · source vision: Goals-First Kustomer (May 2026, Aditya Subramaniam) · baseline: current
Kustomer production navigation.