

Milestone 1: Goals-based AI Monitoring

Goals-Based AI Monitoring

Concept: AI Monitor & Goals (Revised)

Aligned with [AI Monitoring & Org-Level Metrics Tech Spec](#) (April 2026)

Author	Aditya Subramaniam
Status	Draft (Revised)
Productboard	AI Monitor, AI Computed Fields, Goals & Reporting
Linear	TBD

Team

Product	Aditya Subramaniam
Engineering	Sylvia Lujo Kevin Lee Michael Meli
Design	Yvonne Doll

Consulted / Reviewed

TBD — to be populated after initial review cycle.

Changelog

Initial draft created. April 17, 2026

Revised to align with AI Monitoring & Org-Level Metrics tech spec. April 27, 2026

Business Goals

Improve admin confidence in AI automation performance by providing continuous, AI-evaluated monitoring of every AI-handled conversation, with configurable goals that track performance against measurable targets over time. This milestone introduces a closed-loop monitoring system where Quality Monitors score live conversations, Suggestions recommend improvements, Goals track business outcomes, and Anomaly Detection surfaces runtime behavioral drift.

Key Results

- Increase admin adoption of AI-powered performance monitoring (tracked via Goals, Quality Monitors, and Computed Field feature usage).
- Reduce mean time to detect AI performance degradation from days to hours.
- Achieve 80%+ admin satisfaction with AI performance visibility (measured via in-app feedback).
- Establish baseline AI Computed Field coverage across 100% of AI-handled conversations within 30 days of deployment.

Concept Summary

What we're building

Computed Fields — A new field type on the Conversation, Customer, and Company Classes whose value is produced by an LLM-based or deterministic tool. Each Computed Field has its own scope, event trigger, computation (via a tool config), and output configuration. Three default Computed Fields ship out of the box: AI-Generated CSAT (conversation-scoped, 1–5 scale), Customer Health Score (customer-scoped, composite of recent sentiment, resolution patterns, and escalation rate), and Company Health Score (company-scoped, aggregated at the company level). A 3-per-Class limit is enforced. Fields are available across the platform for routing, reporting, and automation. These fields are not seeded historically upon enable.

Quality Monitors (under Analyze) — Post-deployment evaluations that score live, closed conversations against explicit, configurable criteria. Each Quality Monitor defines multiple weighted criteria (each with its own rating categories and pass/fail determination), a pass threshold, segmentation conditions controlling which conversations are evaluated, a context profile defining what data the LLM judge receives, and an optional suggestion configuration. Two default monitors ship: an AIC AI-CSAT Monitor (if AIC is enabled) and an AIR Copilot Usage Monitor (if AIR is enabled). One additional custom monitor is permitted in M1, with room to expand as business need arises. Quality Monitors cover both AI-handled and human-assisted conversations via the Kind field ("aic", "air", or "all").

Suggestions — AI-proposed improvements generated in response to Quality Monitor failures or Anomaly Detection events. Each Suggestion targets a specific resource (a procedure, KB article, tool config, or agent), includes a proposed change with before/after content, and carries a priority (high, medium, low). Suggestions bridge the gap between observing a problem and taking action — not just showing that quality has degraded, but proposing what to change next. Admins review and apply or dismiss suggestions manually.

Anomaly Detection — Extends the existing AIC 2.0 Validation Agent into a customer-visible product surface. The Validation Agent already detects runtime deviations such as procedure non-adherence, KB non-adherence, and tool misuse. This milestone turns those internal signals into reportable anomaly events visible through a dedicated Anomaly Detection Report, and feeds them into the Suggestion pipeline so detected drift can lead to concrete improvement proposals.

Goals — Configurable performance targets tied to Computed Fields. Each goal defines a field reference, a target value, a direction (higher is better or lower is better), and a unit (percent, count, score, or currency). Admins can use three default goals (Increase CSAT, Improve Customer Sentiment, Improve Net Retention) or create custom goals referencing any Computed Field. Goals can optionally be scoped to a specific AI Agent via a Goal Resource

reference. Goals are used strictly for monitoring and alerting in this milestone, not for building or configuring AI agents.

Goal-Level Performance Reporting — A reporting surface that shows progress against each configured goal. Goal progress is tracked via point-in-time snapshots (GoalProgressSnapshot events) written each time a qualifying conversation is evaluated, capturing the field value, whether the linked evaluation passed, and the conversation context. Reporting surfaces these snapshots as trend data over time, broken down by procedure and topic where applicable, with flagged conversations shown alongside per-field scores.

Alerting — Configurable alert thresholds on Quality Monitors and Goals. Admins can opt into alerts for specific monitors and configure delivery cadence (hourly or daily). Alerts are delivered via email and in-app notifications. Alert delivery is coalesced to one notification per monitor per delivery window to prevent notification spam.

What we're not building

- Using goals to build, configure, or optimize AI agents. Goals are monitoring-only in this milestone.
- Automated remediation. When a goal is missed or a monitor fails, we alert the admin and may generate a Suggestion, but we do not automatically adjust procedures, routing, or agent behavior.
- AI Computed Fields admin UI for standalone field management. Fields are created implicitly when configuring monitors and goals. A dedicated field management UI is deferred to a future milestone.
- Custom aggregation strategies for Customer and Company level fields. M1 uses a default rolling average. Configurable aggregation (weighted average, percentile) is deferred.
- Goal-linked custom reporting queries. Tracking higher-order business outcomes such as subscription save rate, average order value, or revenue via saved reporting queries is deferred to M3.
- Monitoring all conversations. Quality Monitors evaluate AI-handled and AIR-assisted conversations. Extending to fully unassisted human conversations is planned for a future milestone.
- Kobject-scoped Computed Fields. Kobject fact pipelines require Debezium connectors and new dbt models and are deferred.

Milestone Phasing

This concept aligns with the four-phase delivery model defined in the tech spec. The concept doc focuses on the M1 customer experience but notes M0 dependencies and M2/M3 extensions for context.

M0 – Internal Foundation (Target: May 8)

- AI-Generated CSAT, Customer Health Score, Company Health Score, and general Computed Fields available internally.
- Quality Monitors, Suggestions, Anomaly Detection, and Computed Fields available internally behind internal-only controls with reasonable per-org safety caps.
- Basic cost and usage instrumentation so internal dogfooding can inform the external rollout plan.

M1 – Starter Experience + Customer Beta (Target: June 2)

- Every org that enables Kustomer AI sees a working quality monitoring experience within one click, without authoring anything themselves.
- Three pre-built Computed Fields visible to customers: AI-Generated CSAT, Customer Health Score, and Company Health Score, each with an out-of-the-box report on the reporting page.
- Beta for Quality Monitors (2 OOTB + 1 custom allowed), Suggestions, Goals, and Anomaly Detection Reports.
- Alerting available for opted-in monitors with hourly or daily cadence via email and in-app notifications.
- A “Kustomer AI Monitoring” dashboard landing page that surfaces these defaults.

M2 – Customer Authoring + Billing Limits (Target: June 15)

- Customers can author their own Computed Fields and Quality Monitors, subject to billing limits per tier/contract.
- Usage accounting and limit enforcement surface the relevant execution and cost constraints.

M3 – Goal-Linked Custom Reporting Queries (Target: July 6)

- Customers can link a custom reporting query to their goal, allowing them to track higher-level objectives such as subscription save rate, revenue, and average order value.
- Goals authored against fieldRef in M1/M2 remain backward-compatible. The integration surface is additive.

Problems & Solutions

Business Challenges

Admins deploying AI automations today lack continuous, AI-evaluated metrics. They can see conversation counts and basic resolution rates, but cannot track customer satisfaction, sentiment trends, or churn risk at the conversation and customer level.

Solution: Introduce Computed Fields and Quality Monitors to continuously evaluate every AI-handled conversation. Computed Fields write AI-derived scores (CSAT, sentiment, churn risk) onto Conversation, Customer, and Company Classes as first-class attributes available across routing, reporting, and automation. Quality Monitors score conversations against multi-criterion rubrics with weighted pass/fail evaluation. Default monitors and Computed Fields are provided out of the box so admins see value immediately upon deployment.

Admins cannot set measurable targets for AI performance or detect degradation before it impacts customers.

Solution: Goals allow admins to define a target value and direction for any Computed Field. Alerting notifies admins when a Quality Monitor detects degradation at a configured cadence (hourly or daily). The reporting surface shows real-time progress against each goal, making it easy to spot trends and take action before metrics deteriorate further.

When problems are detected, admins have no systematic path from observation to action.

Solution: Suggestions are auto-generated when Quality Monitors fail or Anomaly Detection events fire, proposing concrete changes to procedures, KB articles, tool configs, or agent settings. Each suggestion includes before/after content and a rationale, so admins can review and apply fixes without starting from scratch.

Technical Challenges

Computed Fields must be evaluated on every AI-handled conversation without meaningfully impacting conversation latency or throughput.

Solution: Evaluate asynchronously post-conversation so the customer experience is never impacted. Implement queue-based processing with backpressure handling and debouncing to prevent redundant evaluations from automations that repeatedly open/close conversations.

Quality Monitor results generate high data volumes that must be retained for both operational access and long-term analytics.

Solution: Dual-store architecture. MongoDB retains operational data with a 30-day TTL on evaluation results (conversationevaluationresults collection) for recent drill-down access. ClickHouse retains analytical data for 24+ months via staging, intermediate, and mart

models. After 30 days, the UI falls back to ClickHouse summary data (evaluation outcome, score, pass/fail) without per-criterion rationale detail.

Jobs To Be Done

ID	PERSO NA	JOB	SUCCESS INDICATOR
JOB 1	Admin	When deploying an AI automation, I want to define measurable performance goals, so that I can track whether my AI is meeting my standards over time.	Admin can configure goals tied to Computed Fields with target values, directions, and units, and can do so using defaults or custom goals, without needing to write queries or build reports manually.
JOB 2	Admin	When monitoring AI performance, I want AI to automatically score every AI-handled conversation on metrics like CSAT, sentiment, and churn risk, so that I have continuous, objective data without needing to manually review conversations.	Quality Monitors evaluate every qualifying conversation and write scores to Computed Fields on the Conversation and Customer Classes automatically without admin intervention after initial configuration.
JOB 3	Admin	When a performance metric degrades, I want to be alerted at a cadence I choose, so that I can investigate and fix issues before they impact more customers.	Admin receives alerts via email or in-app at hourly or daily cadence when an opted-in Quality Monitor detects degradation, without needing to manually check the dashboard.
JOB 4	Admin	When reviewing AI performance, I want to see goal-level reporting broken down by procedure, topic, and time period, so that I can pinpoint exactly where my AI is underperforming.	Admin can view reporting with per-goal trend data, procedure and topic breakdowns, and flagged conversations from a single view without switching between multiple reports.
JOB 5	Admin	When investigating flagged conversations, I want to see the per-criterion Quality Monitor scores alongside the reason it was flagged, so that I can quickly understand what went wrong.	Admin can see per-criterion scores, weights, pass/fail status, and rationale for each conversation with a single click to view the full conversation.
JOB 6	Admin	When a Quality Monitor flags a failure, I want an AI-generated Suggestion that proposes a concrete fix, so that I can act on problems instead of just observing them.	Admin receives Suggestions targeting a specific procedure, KB article, or tool config with a proposed change (before/after), rationale, and priority, and can apply or dismiss from the Suggestions view.
JOB 7	Admin	When my AI agent deviates from expected behavior at runtime, I want to see those	Anomaly Detection events from the AIC 2.0 Validation Agent are visible in

		anomalies surfaced in a report and linked to improvement suggestions, so that I can catch drift before it becomes a pattern.	a dedicated Anomaly Detection Report and feed into the Suggestion pipeline.
JOB 8	Admin	When I need to track a custom metric specific to my business, I want to add a new goal tied to a Computed Field, so that I can monitor what matters most to my organization.	Admin can create custom goals by specifying a Computed Field reference, target, direction, and unit from the Goals configuration screen without needing engineering support.
JOB 9	Admin	When configuring Quality Monitors, I want to define weighted criteria with rating categories and a pass threshold, so that the scoring matches my business context and priorities.	Admin can add, weight, and mark criteria as essential, configure rating categories per criterion, and set a pass threshold from the Quality Monitor configuration screen.

Initial Spec Ideas

Exploratory first thoughts to pressure-test feasibility and direction. Non-binding and subject to change after review.

Computed Fields

Computed Fields are AI-generated or deterministic attributes written back onto Kustomer resources (conversations, customers, companies). They are standalone product primitives with their own scope, event trigger, computation pipeline, and output configuration.

Default Computed Fields

- **AI-Generated CSAT** — Conversation-scoped. LLM-inferred CSAT from a closed conversation's transcript, written to a conversation custom field on a 1.0–5.0 scale.
- **Customer Health Score** — Customer-scoped. LLM-graded composite of recent conversation sentiment, resolution patterns, and escalation rate over the last 90 days. Rolling average on the Customer Klass.
- **Company Health Score** — Company-scoped. Same pattern aggregated at the company level.

Computed Field Configuration

- **Name** — Friendly label for the field.
- **Description** — What the field measures and why it matters.
- **Scope** — Which resource type the field lives on: Conversation, Customer, Company, or Kobject (Kobject deferred).
- **Trigger** — The event that initiates computation (e.g., `kustomer.conversation.update:done`).
- **Computation** — References a tool config that defines how the field is calculated (LLM-based or deterministic).
- **Output** — The custom field key the value is written to, the value type, and an optional rationale field name for AI explanations.

Limits and Constraints

- 3 Computed Fields per Klass (each default counts toward the limit).
- Fields are not seeded historically upon enable.
- Evaluation is asynchronous post-conversation (fire-and-forget with observability).
- Debouncing: if multiple trigger events fire for the same resource within a short window, only the last is processed.
- Optimistic locking with PATCH operations on subjects, retrying on conflict.

Quality Monitors (Analyze > Quality Monitors)

Quality Monitors are post-deployment evaluations that score live, closed conversations against explicit criteria. They are the core operator-facing surface for understanding whether AI-handled and human-assisted conversations are behaving as intended.

Default Monitors

- **AIC AI-CSAT Monitor** (if AIC is enabled) — Evaluates AI-handled conversations for resolution quality, accuracy, tone, and customer satisfaction signals.
- **AIR Copilot Usage Monitor** (if AIR is enabled) — Evaluates human-assisted conversations for copilot adoption, suggestion acceptance, and quality impact.

Monitor Configuration

- **Name / Description** — Friendly label and explanation of what the monitor evaluates.
- **Kind** — Which conversations are evaluated: "aic" (AI-handled), "air" (human-assisted), or "all".
- **Segmentation** — Conditions (structured) and/or natural-language conditions controlling which conversations match.
- **Context Profile** — Controls what data the LLM judge receives: messages, automation config, execution traces, copilot messages, relevant KB articles, customer data, and/or kobject data.
- **Analysis** — The scoring rubric:
 - **Criteria** — Each criterion has a name, type ("ai", "tool", or "condition"), instructions or tool config reference, rating categories (each with a name, description, pass/fail flag, and score percentage), a weight (relative importance), and an essential flag (failing an essential criterion = entire evaluation score becomes 0%).
 - **Pass Threshold** — A 0.0–1.0 threshold that the overall weighted score must meet for the conversation to pass.
- **Output Field** — Optional fields where the overall score and pass/fail result are written as Computed Fields on the conversation.
- **Suggestion Config** — Whether suggestions are enabled for this monitor, and which target types are eligible (procedure, KB article, tool config, agent).
- **Goal IDs** — Links this monitor to one or more Goals for progress tracking.

One Custom Monitor in M1

Beyond the two defaults, admins can create one additional custom monitor in M1. This allows admins to define their own criteria, weights, and thresholds for a use case specific to their business. The limit can be expanded as business need arises.

Suggestions

Suggestions are AI-proposed improvements generated when a Quality Monitor fails or an Anomaly Detection event fires. They are paired with monitors and anomalies rather than positioned as a standalone product surface.

Suggestion Fields

- **Type** — What kind of change: `procedure_edit`, `kb_update`, `tool_config`, or `agent`.
- **Status** — pending, applied, or dismissed.
- **Priority** — high, medium, or low.
- **Target** — The specific resource the suggestion applies to (resource type, ID, and name).
- **Proposed Change** — Change type (edit or create), with before and after content so admins can see exactly what would change.
- **Rationale** — Explanation of why the change is recommended, linked to the Quality Monitor failure or anomaly that triggered it.

Admins review suggestions and apply or dismiss them manually. Dismissed suggestions require a dismiss reason. Repeated suggestions for the same target are deduped within a bounded time window.

Anomaly Detection

Anomaly Detection extends the AIC 2.0 Validation Agent (already shipped in production) into a customer-visible surface. The Validation Agent inspects an agent's in-flight or just-completed reasoning and flags deviations such as procedure non-adherence, KB non-adherence, and tool misuse.

- **Anomaly events** are emitted with an anomaly type, description, trigger source, and affected resource.
- **Anomaly Detection Report** surfaces these events in a dedicated reporting view.
- **Suggestion pipeline integration** — Anomalies feed into the same Suggestion pipeline as Quality Monitor failures, generating proposed fixes when actionable.

Goals (Settings > Goals)

Goals are configurable performance targets that reference Computed Fields. Each goal defines what metric to track, a target value, and the direction that indicates improvement.

Default Goals

- **Increase CSAT** — References AI-Generated CSAT Computed Field. Direction: higher is better. Target: configurable (suggested default 4.5).
- **Improve Customer Sentiment** — References Customer Health Score Computed Field. Direction: higher is better.
- **Improve Net Retention** — References a churn risk Computed Field. Direction: lower is better.

Goal Configuration Fields

- **Goal Name** — Short label (e.g., "Improve CSAT").
- **Description** — Natural language objective.
- **Field Refs** — One or more Computed Field paths tracked by this goal.
- **Goal Resource** (optional) — Scopes the goal to a specific resource type and ID (e.g., a specific AI Agent/automation).
- **Metric** — The primary tracked field, consisting of:
 - **Field Ref** — Must match one entry in Field Refs.
 - **Target** — Numeric target value.
 - **Direction** — "higher_is_better" or "lower_is_better".
 - **Unit** — "percent", "count", "score", or "currency".
- **Evaluation IDs** — Links to Quality Monitors whose results contribute to goal progress tracking.

Goal Progress Tracking

Goal progress is tracked via GoalProgressSnapshot events. Each snapshot captures the goal ID, field reference, current field value, the conversation and evaluation that produced it, and whether the evaluation passed. These snapshots are aggregated daily in ClickHouse (avg/min/max values, pass/fail counts) and surfaced in the reporting view as trend data.

Open Questions

- Do we need a maximum number of goals per org, or is this unlimited?
- Should Goals support being scoped to multiple automations, or is one-to-one sufficient for M1?

- How do we handle the case where an admin creates a goal referencing a Computed Field that no monitor writes to? Should we warn them, or allow it and show "no data"?

Alerting

Alerting allows admins to opt into notifications when Quality Monitors detect degradation.

Alert Configuration

- **Opt-in per monitor** – Admins choose which monitors trigger alerts.
- **Cadence** – Hourly or daily.
- **Delivery channels** – Email and in-app notifications.
- **Coalescing** – One notification per monitor per delivery window. Delivery retries are rate-limited to avoid notification spam.

Open Questions

- Should alerts support delivery to specific users or teams, or is org-admin-only sufficient for M1?
- Should goal-level alerts (not just monitor-level) be supported in M1, or deferred?

Reporting Surfaces

The following reporting surfaces ship in M1:

- **AI-Generated CSAT Report** – OOTB report on the reporting page for the AI-CSAT Computed Field.
- **Customer Health Score Report** – OOTB report for the customer-scoped health score.
- **Company Health Score Report** – OOTB report for the company-scoped health score.
- **Anomaly Detection Report** – Surfaces AIC 2.0 Validation Agent anomaly events.
- **Kustomer AI Monitoring Dashboard** – Landing page surfacing default monitors, goals, recent alerts, and flagged conversations.

Goal-level trend data (from GoalProgressSnapshot events) is available through the monitoring dashboard. Detailed procedure-level and topic-level breakdowns depend on the data available in ClickHouse intermediate models (evaluation performance hourly, goal progress daily).

Data Retention

Quality Monitor results are stored in MongoDB with a 30-day TTL on the conversation evaluation results collection. For conversations older than 30 days, the UI falls back to

ClickHouse summary data (evaluation outcome, score, pass/fail) without per-criterion rationale detail. ClickHouse retains analytical data for 24+ months. Suggestions retain a 30-day TTL past resolution; pending suggestions never expire via TTL.

Existing Screens (Preserved)

The following existing screens are preserved as-is within the AI for Customers navigation and are not modified in this milestone:

- **Build** – Guidance, Knowledge Sources, and Procedures configuration.
- **Test** – Test categories and evaluation scores (pre-deployment Evaluations, distinct from Quality Monitors).
- **Deploy** – Conditions, Smart Routing, and Review Evaluations.
- **Settings > General** – AI Automation Settings, Guardrails, Handoff, Conversation Start, Verification, Abandonment, Fallback, and Business Hours.
- **Analyze > Observability** – Traces and execution logs.

Out of Scope

- **Goal-driven AI agent optimization:** Using goals to automatically adjust AI agent behavior, procedures, or routing. Goals are monitoring-only in M1.
- **Automated remediation:** When a goal is missed, we alert the admin and may generate a Suggestion, but do not automatically fix the issue.
- **AI Computed Fields admin UI:** A dedicated screen for managing (creating, renaming, deleting) Computed Fields is deferred. In M1, fields are created implicitly via monitor and goal configuration.
- **Custom aggregation strategies:** Configurable aggregation for Customer and Company level fields is deferred. M1 uses a default rolling average.
- **Goal-linked custom reporting queries:** Tracking higher-order business outcomes via saved reporting queries is deferred to M3. Goals in M1/M2 use fieldRef; M3 will add a queryId or similar reference additively.
- **Kobject-scoped Computed Fields:** Requires Debezium connector and new dbt models. Deferred.
- **Historical backfill:** Computed Fields are not seeded historically upon enable.
- **User-authored raw SQL:** Authoring happens through structured config, not free-form expressions.

Risks & Mitigation

Evaluation latency and throughput: Running AI evaluations on every conversation could introduce latency or create a processing bottleneck at scale.

Mitigation: Evaluate asynchronously post-conversation. Implement debouncing so repeated trigger events for the same resource within a short window only process the last one. Start with the default monitors and measure throughput before expanding custom monitor limits.

LLM cost runaway: One enterprise org enabling several Computed Fields and monitors across high-volume traffic could cost far more than expected.

Mitigation: Per-org daily execution caps. 3-per-Klass Computed Field limit. Monitor limits (2 OOTB + 1 custom in M1). Cost alerting when org crosses thresholds. Billing-based enforcement before broad M2 rollout.

Concurrent automation conflicts: Computed Fields introduce another automation system running alongside Workflows, Business Rules, and Queues/Routing, increasing the risk of concurrent writes to the same resource.

Mitigation: Optimistic locking with retry on conflict for Computed Field PATCH operations. Batching multiple Computed Fields for the same Klass and trigger into a single PATCH where possible.

AI Computed Field accuracy: AI-generated scores may not always align with ground truth, especially for nuanced conversations.

Mitigation: Allow admins to customize evaluation criteria and prompts. Surface flagged conversations with full per-criterion scores so admins can audit AI evaluations. Consider adding a human feedback loop (thumbs up/down on AI scores) in a future milestone.

30-day TTL impact on drill-down: After MongoDB TTL expiration, per-criterion rationale detail is unavailable.

Mitigation: Fall back to ClickHouse summary data (outcome, score, pass/fail). Document this limitation clearly. Evaluate whether mirroring rationale to a cheaper store is warranted based on customer feedback.

Alert fatigue: Poorly tuned monitors or repeated failures could produce noisy notifications.

Mitigation: Opt-in alerts only. One notification per monitor per delivery window. Delivery throttles. Easy mute/disable controls.

Scope creep into agent optimization: Stakeholders may push to use goals for automated agent tuning in M1.

Mitigation: Clearly scope M1 as monitoring-only. Goals power dashboards, alerts, and suggestions, not agent behavior. Document the vision for goal-driven optimization so stakeholders know it's planned, just not in this milestone.

Competitor Analysis

See the separate [Competitive Research document](#) for a full landscape analysis. Key areas evaluated for this initiative include how competitors (Intercom, Zendesk, Sierra, Decagon, Siena, Gorgias) surface AI performance metrics to admins, whether they offer configurable goals or targets, how they handle AI-evaluated quality scoring, and what alerting mechanisms exist for AI performance degradation. The core finding remains: no competitor has shipped a true goals-first configuration experience where admins declare business outcomes and the platform auto-steers monitoring, evals, and reporting. Quality Monitors, Computed Fields, and Goals together position Kustomer to claim this whitespace.

-

QA Competitors

Competitive Research: Traditional CX QA Players vs. Kustomer Quality Monitors

Executive Summary

The traditional QA category is undergoing a forced reinvention. Sampling-based, manual QA is dead, GenAI Auto-QA has become table stakes (every vendor on this list claims "100% conversation coverage"), and the new battleground is whether a QA tool can evaluate **AI agents** as rigorously as it evaluates humans. Kustomer's Quality Monitors enters a market with five strategically meaningful incumbents: **MaestroQA, Level AI, Observe.AI, Scorebuddy, and NICE CXone QM (which now absorbs Playvox post the October 1, 2024 acquisition)**.

Key findings about the QA category landscape:

1. **The category has bifurcated into "QA-as-a-feature" inside CCaaS suites (NICE Mpower QM, Salesforce Agentforce QA, Zendesk QA) versus standalone "best-of-breed" QA platforms (MaestroQA, Scorebuddy, Level AI, Observe.AI).** The standalones are increasingly squeezed – NICE's \$249/agent/month Mpower Ultimate Suite explicitly bundles QM, AI-routing, Enlighten Copilot and 1,000+ pre-built AI models, making standalone QA harder to justify on TCO alone.
2. **AI-agent evaluation is the new frontier and most QA incumbents are visibly behind.** Observe.AI is the only vendor with a discretely-marketed "AI Agent Evaluation" product (covering 100% of AI-led conversations for accuracy, empathy, compliance, resolution). MaestroQA has Chatbot QA + integrations for Sierra, Decagon, Forethought, and Agentforce. Scorebuddy claims human + AI scoring. Level AI added a "Unified Intelligence Loop" only in February 2026. NICE QM's 26.1 Auto Score (GA Feb 2026) still primarily targets human agents.
3. **Pricing models are colliding.** Per-agent pricing (MaestroQA \$15+/agent/mo, NICE Mpower \$135–\$249/agent/mo, Scorebuddy seat-based bundles) is being attacked by upstarts pushing per-conversation/per-ticket pricing (\$0.08/audited ticket from Score AI, etc.). The vendors evolving fastest publish hybrid models that decouple AI scoring credits from human evaluator seats.
4. **Goals-based configuration is a market gap.** Across all five vendors, "scorecards with weighted criteria" exist universally but none expose **configurable goals tied to computed fields with goal-level performance reporting and alerting** as a

first-class object. They sell "scorecards" and "metrics," not "goals" with thresholds, owners, and drift detection.

5. **Anomaly detection / drift detection of live AI agents is essentially absent.** None of the five publicly market a runtime "validation agent" that detects AI drift in production. This is Kustomer's biggest white space.

Biggest threats to Kustomer:

- **Observe.AI's** explicit "evaluate humans + AI" positioning, \$44.2M in 2024 revenue, "over 350 successful customer deployments" (per BusinessWire, August 21, 2024), and March 2025 VoiceAI launch make it the most credible AI-era QA player.
- **MaestroQA's** deep helpdesk integrations (Zendesk, Salesforce, Kustomer, Sierra, Decagon, Agentforce, Forethought) and its March 2025 AI Platform launch (LLM Classifiers + AskAI on Amazon Bedrock) make it the most likely overlay on top of Kustomer for customers who want a QA tool that travels across helpdesks.
- **NICE Mpower QM Auto Score** (GA February 2026 in 26.1 release) — by sheer distribution and 11 consecutive years as a Leader in the Gartner CCaaS Magic Quadrant (per NICE/BusinessWire, September 10, 2025) — is the platform play that mid-market and enterprise CX leaders default-evaluate.

Kustomer's differentiation path: Win by being the **only platform that natively unifies the conversation system of record + AI agent (AIC) + agent assist (AIR) + Quality Monitors + Anomaly Detection + Goals**. The standalone QA vendors must integrate; Kustomer just *is*. Lean hard into goals-based config, AIC 2.0 Validation Agent drift detection, and Computed Fields as cross-object KPIs (Customer Health, Company Health, AI-Generated CSAT) — none of the five competitors expose this primitive.

Top 5 strategic recommendations:

1. **Lead with "post-deployment safety net for AI agents"** — explicit Quality Monitors coverage for both AIC (AI-handled) and AIR (human-assisted) is a real wedge against Scorebuddy and Level AI which only recently added AI-agent evaluation.
2. **Price disruption: bundle Quality Monitors into the Kustomer platform** at near-zero incremental cost so customers compare *Kustomer-included QA* vs. **\$15–\$60/agent/month** standalone QA stacks.
3. **Productize Goals as a first-class object** with thresholds, owners, alerting, and goal-level reporting — none of the five competitors have this.
4. **Showcase Anomaly Detection from the AIC 2.0 Validation Agent** as a category-creating capability — call it *runtime QA*, distinct from post-call QA.
5. **Defend Kustomer's installed base from MaestroQA overlay** by tightening native integration: ensure Computed Fields and Quality Monitors expose API parity that any

third-party QA vendor would otherwise require — make Maestro's pitch redundant inside Kustomer.

Landscape at a glance

Vendor	Goals-first config	Eval/sim rigor	Governance depth	2026 pricing signal	Threat level
MaestroQA	✗ Scorecards + AI Classifiers; no native "Goals" object	★★★★★ LLM Classifiers, AskAI, AI Platform with side-by-side prompt grading on Amazon Bedrock	★★★★★ SOC 2 Type 2, HIPAA; auditor-friendly side-by-side review	\$15–\$19/agent/month list; Vendr: \$6K–\$200K+ ARR by team size	High — overlay risk on Kustomer
Level AI	✗ Scorecards via QA-GPT; no goals	★★★★★ QA-GPT, claims 90%+ scorecard accuracy, sandboxed eval	★★★★★ SOC 2, ISO 27001, PCI, GDPR, Gartner Cool Vendor 2023	Quote-only enterprise; 2023 revenue \$23.4M (Latka); \$73.1M raised	Medium-High — enterprise displacement
Observe.AI	✗ Scorecards/rubrics; no goals object	★★★★★ ★ Auto QA on humans + AI; explicit AI Agent Evaluation product	★★★★★ SOC 2 Type II, HIPAA, GDPR, full audit trails, redaction	Quote-only; \$44.2M 2024 revenue; \$214M total raised	High — strongest AI-eval narrative

Scorebuddy	✗ Customizable scorecards only	★★★★ GenAI Auto Scoring + AI Knowledge (RAG, 2025)	★★★★★ ISO 27001 (UKAS-accr edited), enterprise BI, multi-region	Bundles: Foundation/Accelerate/Elite; ~\$80/agent/mo seat-based reported	Medium – Intercom partner advantage
NICE CXone QM (incl. Playvox)	⚠ Mpower Quality Plans, evaluation forms; not "goals" as a primitive	★★★★★ QM Auto Score (26.1 GA Feb 2026), Enlighten LLMs, 1,000+ pre-built AI models	★★★★★ Gartner CCaaS MQ Leader 11 consecutive years (Sept 2025), FedRAMP, every major regulated certification	\$135 (Essential w/ QM)–\$249 (Mpower Ultimate)/agent/mo list	Very High – bundled platform threat

1. MaestroQA

Product capabilities

Core primitives and admin model. MaestroQA's primitives are scorecards (manual + auto), AI Classifiers (LLM-based metrics), AskAI prompts, RCA dashboards, and coaching workflows. Its AI Platform (March 2025, integrated with Amazon Bedrock) lets QA teams write a prompt, test it side-by-side against real tickets, refine, and then "launch as a trackable metric across 100% of tickets." This is the closest competitor to Kustomer's Computed Fields – but it's metric-as-prompt, not first-class goals tied to thresholds.

Testing/evaluation capabilities. AutoQA combines LLM-based scoring with Phrase Match and process-based logic. Calibration workflows and dispute features exist. Predictive CSAT was added as a playbook in 2025, using LLM Classifiers to score resolution, sentiment, and effort across 100% of conversations – a near-direct overlap with Kustomer's AI-Generated CSAT Computed Field.

Governance features. SOC 2 Type 2, HIPAA-compliant. Side-by-side AI grading with auditor-friendly evidence and reasoning. No public anomaly detection or drift monitoring.

Recent product launches (2024–2026):

- **March 2025: AI Platform / Amazon Bedrock integration** — AskAI, LLM Classifiers, prompt-to-metric workflow.
- **2024: Front, Sierra, Decagon, Agentforce, Forethought integrations** — explicitly to QA AI bot conversations.
- **Predictive CSAT playbook** (2025) — Kustomer Computed Fields lookalike.

Coverage of AI-handled vs human-handled. Strong on both. Dedicated **Chatbot QA** use case page; Upwork case study reduced AI hallucinations using MaestroQA's Chatbot QA framework; Bombas case study (Gladly chatbot QA); Sierra and Agentforce integrations specifically for AI-agent QA.

Auto-QA capabilities. Full LLM-powered AutoQA with custom rubrics; Predictive CSAT; AI Classifiers turning prompts into KPIs.

Integrations. Zendesk, Salesforce Service Cloud, **Kustomer**, Intercom, Front, Aircall, Talkdesk, 8x8, Five9, Amazon Connect, Dialpad, Sierra, Agentforce, Forethought, Decagon, Snowflake. This is the broadest helpdesk integration footprint in the standalone QA category — and yes, MaestroQA already integrates with Kustomer.

Goals-first gap

MaestroQA has **no native "Goals" object**. It has scorecards, AI Classifiers, KPI dashboards, and "performance metrics" — but a QA Manager cannot set "Customer Health Score ≥ 75 by Q2, owner = Jane, alert if drifts $>10\%$ " as a configurable goal. KPIs live in dashboards and AskAI worksheets, not as governed targets with state machines.

SWOT

Strengths. Broadest helpdesk integration footprint; strong customer logos (Etsy, Mailchimp, Peloton, DraftKings, Stitch Fix, Postmates, Lyft, Squarespace, ClassPass, Tails.com, Dialpad, Handy, Scribd); \$31.5M total raised; Forbes (September 29, 2021) headline confirmed "Already Profitable MaestroQA Raises \$25 Million" at Series A (no more recent public profitability update); LLM-native AI Platform launched March 2025; demonstrably extends to AI agents (Upwork chatbot, Bombas, Sierra integration).

Weaknesses. No goals-based primitive; no anomaly/drift detection; no native CRM/conversation system of record; pricing (\$15–\$19/agent/month and \$35K–\$70K mid-market ARR) becomes redundant when bundled QM is included with the underlying CCaaS/CRM platform; mobile-only product is limited.

Opportunities for Kustomer. (1) MaestroQA already integrates with Kustomer — meaning Kustomer's installed base is *the* lowest-resistance market for MaestroQA overlay sales. Kustomer must close that loop with native Quality Monitors before Maestro entrenches. (2)

Maestro's lack of a goals object is a clean differentiation path. (3) Maestro can't see runtime AI-agent drift inside Kustomer's AIC platform — Anomaly Detection is a moat.

Threats to Kustomer. Maestro is the most likely overlay product Kustomer customers buy *in addition to* Kustomer. Their March 2025 AI Platform launch is genuinely strong; their Sierra/Agentforce/Decagon integrations show AI-agent QA ambition. Maestro could position itself as "the QA tool that travels with you across Zendesk + Kustomer + Salesforce."

Pricing

- **Starting list:** \$15/agent/month for QA for Monitoring (Basic); \$19/agent/month per SourceForge.
 - **Vendr benchmark transaction data:** Small (5–25 agents): \$6K–\$18K/year; Mid (25–100 agents): \$18K–\$60K/year; Enterprise (100+): \$60K–\$200K+/year.
 - **Stacked mid-market example (50-agent team, Professional plan):**
~\$35K–\$70K/year — typically \$1,200/agent/year all-in including basic integrations. With Auto QA premium tier and Salesforce/Snowflake integrations, expect the high end.
-

2. Level AI

Product capabilities

Core primitives and admin model. Level AI is positioned as an **LLM-native CX intelligence platform**, with scorecards driven by **QA-GPT** (proprietary LLM trained on contact-center data). Primitives include scorecards, agent screen recording, AgentGPT (agent assist trained on enterprise knowledge), iCSAT (Inferred CSAT), Voice of the Customer, and a new AI Virtual Agent (Feb 2026 expansion).

Testing/evaluation capabilities. QA-GPT claims to cover 85%+ of scorecard questions out of the box and 90%+ accuracy versus a consensus of human QA managers. The platform offers a sandboxed "test, tune, validate" environment for Auto-QA accuracy.

Governance features. SOC 2, ISO 27001, PCI, GDPR. Named a 2023 Gartner Cool Vendor.

Recent product launches (2024–2026):

- **AgentGPT** (Feb 2023) — generative AI agent assist trained on customer-specific data.
- **QA-GPT** — LLM-powered Auto-QA with explainable scoring evidence.

- **January 22, 2025: Inferred CSAT (iCSAT)** – automated CSAT scoring across 100% of interactions (direct overlap with Kustomer's AI-Generated CSAT Computed Field).
- **July 23, 2024: \$39.4M Series C** led by Adams Street Partners with participation from Cross Creek, BrightLoop, and existing investors Battery Ventures and Eniac Ventures (per TechCrunch and Level AI's own announcement); total raised \$73.1M.
- **February 2026: AI Virtual Agent expansion** – Level AI markets a "Unified Intelligence Loop" using "a single set of human-grade benchmarks to continuously evaluate and retrain both virtual and human teams in tandem" – i.e., catching up to Observe.AI on dual human+AI evaluation.

Coverage of AI-handled vs human-handled. Historically strong on humans (Auto-QA, agent screen recording). Only added explicit AI-agent evaluation in February 2026 with the Unified Intelligence Loop. Plays catch-up to Observe.AI here.

Auto-QA capabilities. QA-GPT is genuinely best-in-class for human evaluation: explainable, rubric-based, evidence-grounded, with side-by-side suggested answers. Stated 90%+ scorecard accuracy.

Integrations. Salesforce (via AppExchange "Level AI QA Assist"), CRM/CCaaS/WFM integrations. Less broad than MaestroQA's footprint.

Goals-first gap

Level AI has scorecards and KPI correlation dashboards but **no goals primitive**. iCSAT is a metric, not a goal with thresholds, owners, and alerts. Customers report KPI correlation as a strength, but governance objects do not exist.

SWOT

Strengths. LLM-native architecture and proprietary QA-GPT with high accuracy; named customers Affirm, Brex, Carta, Penske, Vista, Toast, OpenTable, ezCater, MoneyLion, Swiss Re, Topcon, Quinstreet, Bakkt, Globalface; 25%/45%/90% reported customer outcomes (CSAT, agent satisfaction, QA time savings). Healthy funding momentum: \$39.4M Series C in July 2024.

Weaknesses. Quote-only pricing typical of enterprise sales (2–6 week setup per Level AI's own materials); broader CX platform scope means longer implementations; AI-handled conversation evaluation only added Feb 2026; no goals object; no anomaly detection; revenue scale modest – Latka reports \$23.4M in 2023 (most recent named-source figure) versus larger competitors.

Opportunities for Kustomer. Level AI sells the broader CX platform vision; Kustomer can position as "QA that's already inside the CRM you bought" with no integration overhead.

Kustomer's AIC + AIR taxonomy maps naturally to Level AI's "human + virtual agent" loop but already exists natively.

Threats to Kustomer. Level AI is the most well-funded enterprise QA-focused player (\$73.1M raised, Adams Street Partners leading) and is increasingly positioning as a CX operating layer. For mid-market and enterprise contact centers with regulated workloads, Level AI's certifications (SOC 2, ISO 27001, PCI, GDPR) and Gartner Cool Vendor status give it credibility Kustomer must match.

Pricing

- Quote-only via "Schedule a Demo." No published per-agent pricing.
 - Setup time: 2–6 weeks per Level AI's official materials (and longer in third-party reports).
 - Indirect indicator from third-party Uniphore comparison: ~\$35/agent + ~\$1,500/integration + platform fees is typical of this tier; Level AI is presumed to price comparably or higher given enterprise positioning.
 - **Stacked mid-market example (50-agent team):** Modeling at ~\$50–\$80/agent/month for a full QA + agent assist + iCSAT bundle implies **\$30K–\$48K/year** for 50 agents on QA alone, with implementation services adding \$20K–\$50K upfront.
-

3. Observe.AI

Product capabilities

Core primitives and admin model. Observe.AI's primitives span: VoiceAI Agents, ChatAI Agents, Agent Copilot, Coaching Copilot, Insights Copilot, Conversation Intelligence, Auto QA, Manual QA, Screen Recording, and the explicitly-named **AI Agent Evaluation** product. The platform is built on a 30-billion-parameter contact-center-specific LLM trained on hundreds of millions of contact-center conversations.

Testing/evaluation capabilities. Auto QA evaluates 100% of conversations (both human and AI). The dedicated AI Agent Evaluation product scores AI-led conversations for accuracy, empathy, compliance, and resolution with heatmaps and trend lines. Insights feed directly into AI model tuning (prompt adjustments, flow updates).

Governance features. SOC 2 Type II, HIPAA, GDPR-compliant. Encryption at rest/in transit; audit logs; PCI/PII automatic redaction; complete audit trails. AI agents have grounding/guardrails, real-time alerts, warm escalation. This is the strongest governance posture among the five.

Recent product launches (2024–2026):

- **Series C: \$125M (April 2022) led by SoftBank Vision Fund 2; total \$214M raised.**
- **March 2025: VoiceAI Agents launch** – the platform's pivot to autonomous voice automation.
- **March 16, 2025: Acquisition of Dubdub.ai** (~\$15M deal, Indian text-to-speech startup) for voice cloning and multilingual TTS.
- **November 2025: Named Cool Vendor in 2025 Gartner Cool Vendors in Customer Service and Support Technology report.**
- **AI Agent Evaluation product** – discrete product page covering 100% of AI-led conversations.

Coverage of AI-handled vs human-handled. Strongest in the category. Both human Auto QA (since May 2022) and explicit AI Agent Evaluation product. Marketing claim: "Automatically QA 100% of human and AI interactions to monitor agent performance, quality and compliance."

Auto-QA capabilities. Industry-claimed first adaptive AutoQA solution (May 2022). Five-times-faster evaluation completion via auto-suggest, auto-fill, auto-submit. Drag-and-drop rule builder. Best-in-class calibration.

Integrations. 250+ pre-built connectors to contact center and CRM systems (cloud + on-prem). Reported deployable in as little as 6 weeks.

Goals-first gap

Observe.AI has scorecards, rubrics, and KPI tracking but **no goals primitive**. Its AI Agent Evaluation surfaces "where AI excels or falls short" with risk surfacing – closer to anomaly detection than any other competitor – but still not exposed as configurable goals with thresholds and owners.

SWOT

Strengths. Strongest AI-handled evaluation narrative in the market; explicit AI Agent Evaluation product; "over 350 successful customer deployments" per BusinessWire (August 21, 2024); **\$44.2M 2024 revenue** (up from \$30.1M 2023, per Latka); \$214M total raised including \$125M SoftBank-led Series C; 95%+ adoption rate claim; enterprise customers DoorDash, SoFi, Uber, Accolade, Prudential, Concentrix, Cox Automotive, Included Health, JK Moving, National Debt Relief, DailyPay, Bill.com, Public Storage, 23andMe, Pearson; the same August 21, 2024 BusinessWire release reports a "193% increase in bookings for the strategic and enterprise segments over H124."

Weaknesses. Voice-first DNA means digital-channel + helpdesk integration is less deep than MaestroQA; quote-only enterprise pricing; some G2 reviews note implementation drag and

broken keyword matching at scale (one reviewer cited a multi-year deployment that delivered no value); breadth-versus-depth concerns as it expands into voice agents, IVA, and conversation intelligence simultaneously.

Opportunities for Kustomer. Observe.AI is voice-first and contact-center-shaped; Kustomer's digital-first omnichannel CRM is structurally different. Position Quality Monitors as "designed for digital-first AI-CX, not retrofitted from voice analytics."

Threats to Kustomer. Observe.AI is the single most credible AI-era QA vendor and the most likely competitive bake-off opponent for any Kustomer enterprise prospect that prioritizes explicit AI-agent evaluation. Their Gartner Cool Vendor 2025 status and enterprise customer roster give them air cover Kustomer Quality Monitors must compete against.

Pricing

- Quote-only enterprise. Tiers: VoiceAI Agents, Real-time AI, Enterprise Advanced, Enterprise Unlimited.
 - No published per-agent figures. SaaSwothy reports "Custom plan" with industry-typical \$11/month basic-tier reference (not Observe.AI's actual price; this is the speech analytics category average).
 - Industry positioning suggests \$80–\$150/agent/month all-in for the full AI + Auto QA + Real-time AI bundle for mid-market; enterprise contracts run six to seven figures annually for 1,000+ agent deployments.
 - **Stacked mid-market example (50-agent team, Auto QA + Real-time AI bundle):** Modeled at ~\$60K–\$120K/year, plus 6–12 weeks of implementation and consulting services.
-

4. Scorebuddy

Product capabilities

Core primitives and admin model. Customizable scorecards (touted as the most flexible scorecard builder in the category), QA workflows, calibration module, agent dashboards, advanced reporting, root-cause analysis, integrated Coaching, integrated LMS, conversational analytics, embedded BI, GenAI Auto Scoring, and AI Knowledge (RAG-based, 2025 launch).

Testing/evaluation capabilities. GenAI Auto Scoring claims 90%+ accuracy validated against manual evaluators in real-world enterprise deployments. AI Assist summarizes/scores conversations on demand. Supports human + AI agent scoring across multiple channels.

Governance features. ISO 27001 certified by a UKAS-accredited vendor. Multi-region data residency (EU/US/UK in Elite). Compliance reporting, scorecard locking, native data segregation for cross-border ops.

Recent product launches (2024–2026):

- **2024: GenAI Auto Scoring launch** — the AI scoring engine on credit-based usage.
- **November 25, 2024: €5M investment from Foresight Group** — Scorebuddy's only meaningful funding round; total raised \$5.23M per Crunchbase.
- **February 2025: Intercom partnership announced** — Scorebuddy named the QA backbone for 27,000+ Intercom businesses; Intercom uses Scorebuddy internally.
- **October 2025: AI Knowledge launched** — first native contact-center QA vendor with RAG capabilities, grounding AI scoring in each customer's knowledge base.
- **July 2025: Hired Microsoft veteran Claire O'Rourke as CRO.**

Coverage of AI-handled vs human-handled. Markets dual coverage: "Auto-score human and AI agents across multiple channels." CEO Derek Corcoran explicitly: "A bot might achieve high deflection rates but deliver poor customer experiences. Our platform helps businesses maintain human oversight, flagging problematic interactions for review."

Auto-QA capabilities. GenAI Auto Scoring with credit-based usage; AI Knowledge (RAG); AI Transcription via Deepgram partnership.

Integrations. Kustomer (yes, native), Salesforce Service Cloud, Talkdesk, Five9, Intercom, Zendesk, Aircall, Dixa, Amazon Connect, Genesys. Open API.

Goals-first gap

Scorebuddy has the most sophisticated **scorecard** builder in the category and integrated BI for cross-functional reporting, but **no goals primitive** with thresholds, owners, and configurable alerts tied to Computed Fields.

SWOT

Strengths. Intercom partnership is a genuine moat (27,000+ businesses); Tide testimonial; strong customer roster including Allianz, Oxfam, Boohoo, Capital One, SSE Airtricity, Sanlam, Games24x7, Deriv, ConnectPay, AutoPay, Permanent TSB, Sun Life, British Red Cross, Echelon Group, Ryanair, Beachbody; serves 50,000+ agents across 25+ countries (per their November 2024 funding announcement); €5M Foresight investment; ISO 27001 certified; first-to-market with RAG-based AI Knowledge for QA.

Weaknesses. Smallest funded among the five (\$5.23M total per Crunchbase); 51 employees per LeadIQ (July 2025); seat-based pricing (\$80/agent/month per Score AI's competitive

page) makes them vulnerable to consumption pricing; some Scorebuddy users have publicly complained of inflexibility and reverting to spreadsheets (per evaluagent's competitive site).

Opportunities for Kustomer. Scorebuddy already integrates with Kustomer – meaning Kustomer customers using Scorebuddy can be migrated to native Quality Monitors. The seat-based price ceiling makes them attackable on TCO.

Threats to Kustomer. The Intercom partnership shows Scorebuddy can become the "preferred QA partner" inside platforms – this is the partnership model Kustomer must explicitly *block* by making Quality Monitors a native, included feature.

Pricing

- Three packages: **Foundation, Accelerate, Elite** (all quote-based custom pricing).
 - Foundation: scorecards, AI Assist, calibration, dashboards, advanced reports, peer-to-peer scoring, Silver onboarding.
 - Accelerate: adds AI Auto Scoring (500 monthly AI scores included), BI custom dashboards, multi-lingual UI, in-app coaching, dedicated CSM, CCaaS integrations, Gold onboarding.
 - Elite: adds Conversation Analytics, 1,000 monthly AI scores, AI Voice Transcription, Salesforce integration, Open API, SSO, multi-region, AI Optimization Service, Platinum onboarding.
 - Score AI's competitive page reports: "Starting at \$80/per support agent per month with per-agent pricing."
 - 14-day free trial.
 - **Stacked mid-market example (50-agent team on Accelerate):** ~\$80/agent/month × 50 = **\$48,000/year**, plus AI score credit bundles for usage above 500/month included credits.
-

5. NICE CXone Quality Management (incl. Playvox post-acquisition)

Product capabilities

Core primitives and admin model. NICE CXone Quality Management ships inside the broader NICE CXone Mpower platform. Core primitives: Quality Plans, evaluation forms, calibration, coaching, screen + interaction recording, Mpower QM Auto Score (LLM-based scoring, GA February 2026), Enlighten LLMs, 1,000+ pre-built Enlighten AI models, Enlighten Copilot for supervisors, Mpower AI Studio, Mpower Autopilot.

Playvox integration status. NICE acquired Playvox on **October 1, 2024**. Per NICE CFO Beth Gaspich on the Q3 2024 earnings call (November 14, 2024, via Investing.com): *"the Playvox was a technology tuck-in acquisition with a purchase price around \$40 million."* Playvox's reported annual revenue at acquisition was approximately \$75M per CX Today — meaning NICE acquired ~\$75M of revenue for ~\$40M. Playvox is being **repositioned as "Playvox by NiCE Workforce Engagement Management"** — increasingly framed as **WFM for CRM teams** (Salesforce, Zendesk, Intercom users), distinct from CXone's native QM. No public hard sunset date; competitor Calabrio openly speculates about feature deprecation given product overlap. **Daniel Pagel** was named Playvox CEO in February 2025. Playvox AutoQA (originally based on Prodsight 2022 acquisition, sentiment-scoring across text channels) continues to be marketed but Level AI publicly criticizes it: "Playvox records interactions only from text-based channels (email and chat)... Playvox Auto-QA is limited to sentiment scores... Playvox relies on 1st generation keyword-based technology."

Testing/evaluation capabilities. QM Auto Score (26.1 release, GA February 2026): "Automatically score interactions at scale with LLMs powered by NiCE AI Models. Go beyond traditional AI to understand context, sentiment, and intent." Auto Score targets 100% unbiased evaluation coverage. Evaluation Summaries turn complex data into coaching insights. Calibrations widget shows ±5% variance among evaluators.

Governance features. Per NiCE/BusinessWire (September 10, 2025): *"This marks the eleventh consecutive year that NiCE has been named a Leader in the CCaaS Magic Quadrant and the first time NiCE was positioned furthest and highest on the axes."* FedRAMP-supported; the most enterprise-grade certification posture in the category. WCAG 2.2 Level AA compliance, Section 508, European Accessibility Act compliant.

Recent product launches (2024–2026):

- **October 1, 2024: Playvox acquisition closed** at ~\$40M purchase price (per NICE Q3 2024 earnings).
- **June 17, 2025: CXone Mpower Agents launch** — agentic AI agents that work front/middle/back office.
- **2025: Cognigy acquisition completed** for AI agents (voice/chat).
- **February 5, 2026: 26.1 release with QM Auto Score GA** — fully automated QA inside Mpower.
- **December 10, 2025: NICE CXone Mpower local instance live in South Africa.**
- **2026 Fast Company World's Most Innovative Companies (#11 in Applied AI category).**

Coverage of AI-handled vs human-handled. QM Auto Score (Feb 2026) primarily targets human agent evaluation per the public messaging (focuses on "agent performance" and "the full customer experience"). Cognigy (acquired 2025) provides voice/chat AI agents with grounding/guardrails, but a unified bot-QA product is **not** separately marketed under QM. This is a real gap.

Auto-QA capabilities. QM Auto Score (GA February 2026) is brand-new but distribution-advantaged.

Integrations. Native to CXone Mpower; AWS partnership (Bedrock, Nova LLMs, Amazon Q integration); Snowflake collaboration for enterprise data.

Goals-first gap

NICE CXone QM has Quality Plans (evaluation cadence + assignment) and $\pm 5\%$ variance widgets – closer to "goals" than any other competitor, but still not exposed as **configurable goals tied to Computed Fields with goal-level performance reporting and runtime alerting** as a first-class primitive.

SWOT

Strengths. Distribution moat – 11 consecutive years as a Gartner CCaaS Magic Quadrant Leader (per September 8, 2025 Gartner report); bundled platform with \$135–\$249/agent/month QM included; 1,000+ pre-built Enlighten AI models; AWS, Cognigy, Snowflake partnerships; FedRAMP and major regulated certifications; Playvox + Cognigy acquisitions consolidating the WEM and AI-agent layers; 150+ countries of deployment.

Weaknesses. Standalone QA vendors consistently outpace NICE on speed of innovation; Playvox migration uncertainty creates customer risk; QM Auto Score only GA February 2026 – late to LLM-native AutoQA versus Observe.AI (May 2022), Level AI (QA-GPT), Scorebuddy (GenAI Auto Scoring), and MaestroQA; NICE's pricing tiers consolidate value but force customers into the full Mpower bundle for the best AI; the product surface is so broad that QM specifically often gets less roadmap attention than CCaaS/WEM.

Opportunities for Kustomer. NICE is voice/CCaaS-shaped; Kustomer's digital-first omnichannel positioning is structurally different. NICE Mpower QM Auto Score lacks demonstrated AI-agent evaluation depth – Kustomer's AIC + AIR coverage is a real wedge for AI-CX-native organizations.

Threats to Kustomer. When CX leaders are asked "what's your QA tool?", the answer is increasingly "whatever's in our CCaaS" – and that's NICE Mpower QM by sheer market share. The bundled price point neutralizes standalone QA economics.

Pricing

- **Digital Agent: \$71/agent/month** (30+ digital channels, basic).
- **Voice Agent: \$94/agent/month.**
- **Omnichannel Agent: \$110/agent/month.**
- **Essential Suite: \$135/agent/month** (includes QM workflows, evaluations, calibrations, coaching, appeals, screen recording).

- **Core Suite: \$169/agent/month** (adds WFM, AI forecasting/scheduling).
 - **Complete Suite: \$209/agent/month** (adds Interaction Analytics + VoC).
 - **Mpower Ultimate Suite: \$249/agent/month** (Complete Suite + every Enlighten Copilot/Autopilot/Action, 1,000+ AI models, AI Routing, Knowledge Management).
 - **Stacked mid-market example (50-agent team on Essential Suite with QM):** $\$135 \times 50 \times 12 = \$81,000/\text{year}$ for QM-included CCaaS; on full Mpower Ultimate, **\$149,400/year**. Compare to \$48K/year Scorebuddy or \$35–\$70K MaestroQA: NICE includes the full CCaaS, voice telephony, recording, WEM, *and* QM in the same price.
-

Where Kustomer wins: the strategic playbook

1. Goals as a first-class object – claim the category gap

None of the five competitors expose configurable Goals tied to Computed Fields with thresholds, owners, alerting, and goal-level performance reporting. They sell scorecards and metrics. **Kustomer should productize Goals as the singular primitive that ties every Computed Field (AI-Generated CSAT, Customer Health Score, Company Health Score) to operational accountability.** Marketing positioning: *"Other QA tools score conversations. Kustomer scores your goals."*

2. Lead with Quality Monitors for AI agents (AIC) – the most defensible wedge

Only Observe.AI has a dedicated AI Agent Evaluation product. MaestroQA has integrations (Sierra, Decagon, Forethought, Agentforce). Scorebuddy and Level AI have only recently added the claim. NICE QM Auto Score doesn't market AI-agent evaluation. **Kustomer Quality Monitors covering both AIC (AI-handled) and AIR (human-assisted) conversations natively, on the same conversation system of record, is a category-creating product.** Lead every demo with this.

3. Anomaly Detection from AIC 2.0 Validation Agent – runtime QA versus post-call QA

Every competitor does post-call/post-deployment QA. **None has runtime drift detection of AI agents in production.** Kustomer should rebrand Anomaly Detection as *Runtime QA* and position it as the only safety net that catches AI-agent drift before customers notice. This is the single most differentiated capability in the listed roadmap.

4. Computed Fields as the cross-object KPI engine

MaestroQA has AI Classifiers and AskAI. Scorebuddy has BI. Level AI has iCSAT. None of them have a unified primitive that puts AI-generated attributes on conversations, customers, *and* companies — and feeds them into Goals. **Kustomer's Computed Fields enable Customer Health Scores and Company Health Scores that QA-only tools cannot replicate** because they don't own the conversation/customer/company data model.

5. Bundled, near-zero-marginal-cost pricing — disrupt the per-agent QA stack

Standalone QA pricing stacks \$15–\$80/agent/month *on top of* the underlying CRM/CCaaS spend. NICE solves this by bundling at \$135–\$249/agent/month. **Kustomer should bundle Quality Monitors at near-zero marginal cost** so the buyer's decision becomes: “*buy Kustomer with Quality Monitors included, or buy Kustomer + \$50K–\$150K/year of MaestroQA/Scorebuddy/Level AI.*” Win on TCO simplicity.

Pricing positioning

Vendor	List price (per agent/month)	50-agent stacked annual	Bundled or standalone
Kustomer Quality Monitors	Bundle at near-zero marginal cost (recommended)	Included in Kustomer subscription	Bundled
MaestroQA	\$15–\$19 list; Vendr negotiated \$20–\$60+ effective	\$35K–\$70K (Professional)	Standalone overlay
Level AI	Quote-only; ~\$50–\$80 implied	\$30K–\$48K + implementation	Standalone enterprise
Observe.AI	Quote-only; ~\$80–\$150 implied bundle	\$60K–\$120K+	Standalone enterprise
Scorebuddy	Reported ~\$80 seat-based	\$48K + AI score credits	Standalone with Intercom partnership
NICE CXone Mpower (QM)	\$135 (Essential w/ QM) – \$249 (Ultimate)	\$81K–\$149K (full CCaaS+QM)	Bundled with CCaaS

Kustomer's price disruption only works because Quality Monitors becomes part of the platform value capture, not a separate SKU — defending against MaestroQA/Scorebuddy overlay attacks and against NICE Mpower's bundled platform play simultaneously.

Conclusion

The traditional CX QA category is in transition. Sampling-based manual QA is dead. Five vendors have meaningful strategic relevance to Kustomer Quality Monitors:

- **Observe.AI** is the strongest AI-era QA narrative (\$44.2M revenue, AI Agent Evaluation product, \$214M raised) — the most credible bake-off opponent for AI-CX prospects.
- **MaestroQA** is the most likely overlay product on Kustomer's installed base (broadest helpdesk integrations, March 2025 AI Platform on Amazon Bedrock, Etsy/DraftKings/Mailchimp/Peloton roster).
- **NICE CXone Mpower QM** (with Playvox folded in at a ~\$40M purchase price) is the bundled platform threat — distribution moat at \$135–\$249/agent/month makes standalone QA economically harder.
- **Level AI** is the well-funded enterprise contender (\$73.1M raised, QA-GPT, Affirm/Brex/Carta/Penske roster) catching up on AI-agent evaluation in February 2026.
- **Scorebuddy** is the partner-driven specialist (Intercom partnership, ISO 27001, RAG-based AI Knowledge) that demonstrates how a focused QA play can scale.

Kustomer's path to win is *not* to out-feature them on scorecards. It is to **own the goals primitive, own runtime AI-agent QA, own the Computed Fields KPI engine across conversations/customers/companies**, and **bundle Quality Monitors at near-zero marginal cost** so the buyer's TCO math is structurally tilted in Kustomer's favor.

The standalone QA vendors must integrate; Kustomer just *is*. Make that the entire strategic narrative.

Definitions

Definitions

Computed Fields	Computed Fields are AI-generated or deterministic attributes written back onto Kustomer resources (conversations, customers, companies). They are standalone product primitives with their own scope, event trigger, computation pipeline, and output configuration.
Goals	Goals are configurable performance targets that reference Computed Fields. Each goal defines what metric to track, a target value, and the direction that indicates improvement.
Monitors	Quality Monitors are post-deployment evaluations that score live, closed conversations against explicit criteria. They are the core operator-facing surface for understanding whether AI-handled and human-assisted conversations are behaving as intended.
Suggestions	Suggestions are AI-proposed improvements generated when a Quality Monitor fails or an Anomaly Detection event fires. They are paired with monitors and anomalies rather than positioned as a standalone product surface.
Anomaly Detection	Anomaly Detection extends the AIC 2.0 Validation Agent (already shipped in production) into a customer-visible surface. The Validation Agent inspects an agent's in-flight or just-completed reasoning and flags deviations such as procedure non-adherence, KB non-adherence, and tool misuse.

Direct Kustomer Competitors

Goals-first AI governance: Kustomer's competitive opening

Executive summary

The "goals-first" category is largely unclaimed whitespace – and the window closes fast.

- Of seven AI-CX competitors analyzed, **none has shipped a true goals-first configuration experience** where admins declare business outcomes and the platform auto-steers agent behavior, evals, procedures, and reporting.
- Sierra and Siena come closest in philosophy. Intercom explicitly promised "Goals" at Pioneer 2025 but has not shipped it. Zendesk and Decagon treat outcomes as a pricing story, not a configuration paradigm.
- This is Kustomer's most exploitable strategic gap in the current market.

Three structural themes define the 2026 landscape:

- **Outcome-based pricing is table stakes** (\$0.90–\$2/resolution across Intercom, Zendesk, Sierra, Decagon, Gorgias, Siena) – but buyers struggle with unpredictable bills, double-billing, and ambiguous "resolution" definitions.
- **Eval and simulation rigor is the new enterprise moat:** Sierra runs 35K+ simulated tests/day; Decagon publishes offline → online A/B; Intercom launched Simulations; Zendesk, Gorgias, and DigitalGenius have meaningful gaps.
- **Governance is bifurcating** into transparency-first (Zendesk AI Reasoning, Gorgias) vs. policy-enforcement-first (Sierra's multi-supervisor architecture, Decagon's Watchtower). Enterprise buyers increasingly want both.

Biggest threats to Kustomer:

- **Sierra** – \$635M raised, \$10B valuation, \$150M+ ARR, 40% of Fortune 50. Ghostwriter (March 2026) targets mid-market directly.
- **Intercom Fin** – \$0.99/outcome with announced Goals vision. March 2026 Pro add-on paywall introduces pricing friction Kustomer can exploit.
- **Zendesk** – Forethought acquisition (closed April 2026) adds self-improving AI, specialized agents, and native voice at massive install-base scale.

Kustomer's differentiation path:

- Ship a **true goals-first UX first** — admin selects an outcome (deflection, CSAT, cost reduction, sales assist) and the platform auto-generates starter procedures, guardrails, eval rubrics, and dashboards tuned to that goal. No competitor has this.
- **Include goal-aligned analytics in the core product** rather than paywalling it (directly exploiting Intercom's Pro add-on backlash).
- Offer **outcome-tier bundles** with predictable monthly commitments, solving the #1 buyer complaint across Sierra, Intercom, and Zendesk.

Top five strategic recommendations:

1. **Ship "Goals" as a first-class admin object in AIC** before Intercom Fin 4. Own the category term.
2. **Build goal-scoped evals** — every goal type auto-generates its own graded test battery. Closes Kustomer's largest feature gap vs. Sierra and Decagon.
3. **Launch goal-based pricing bundles** with committed-outcomes model. Undercut on TCO predictability, not per-unit price.
4. **Lead on governance transparency** — ship a hallucination-rate dashboard, configurable confidence thresholds, and PII/toxicity evaluators. No competitor fully delivers this.
5. **Sequence AIR to piggyback on AIC goals** — a rep-side copilot aligned to the same declared goals would be industry-first.

The landscape at a glance

Vendor	Goals-first config	Eval/sim rigor	Governance depth	2026 pricing signal	Threat level
Intercom Fin	Vision only (Pioneer '25)	High (Simulations GA'd)	Strong (AIUC-1 audit)	\$0.99/outcome + seats + Pro \$99/mo+	Very High
Zendesk	None (pricing framing only)	Weak (no native A/B)	Medium (CoT visibility)	\$115–\$209/agent + ~\$1.50–\$2/resolution	Very High
Sierra	Closest in market (Journeys)	Category-leading	Strongest (supervisors, PCI L1)	Outcome-based, ~\$1–\$1.50; \$150K+ ACV	Very High
Decagon	Implicit (AOPs + Experiments)	Very strong (offline→online A/B)	Strong (Watchtower)	Per-conversation or per-resolution; ~\$386K median ACV	High

Vendor	Goals-first config	Eval/sim rigor	Governance depth	2026 pricing signal	Threat level
DigitalGenius	None (use-case driven)	Weak (CSM-led)	Basic	From \$1K/mo + implementation	Low
Siena AI	Explicit ("Set specific goals")	Medium (QA Agent)	Medium-strong	Per-conversation, custom	Medium
Gorgias	None (tone/style knobs)	Weak (preview only)	Basic (no numeric thresholds)	\$0.90-\$1.00/resolution + helpdesk tickets	Medium

Intercom Fin

Product capabilities

Fin's admin model follows a four-stage "Fin Flywheel": **Train** → **Test** → **Deploy** → **Analyze**.

Core primitives:

- **Procedures** (Pioneer 2025) — natural-language + branching logic + code + data connectors for multi-step workflows; supersedes Fin Tasks
- **Guidance** — custom coaching on tone, escalation, brand voice, audience behavior
- **Custom Answers** — discontinued for new customers March 19, 2025
- **Vision** — image interpretation
- **Audiences** — targeting by segment
- 45 languages with real-time translation

Testing:

- **Simulations** (Pioneer 2025, managed availability) — AI-as-customer runs end-to-end Procedure conversations, AI-as-judge grades pass/fail, tests rerun as regression safety nets
- **Batch testing** (up to 50 questions), **Preview** (interactive sandbox), **Controlled Rollout** (deploy to subsets before full launch)

Governance:

- Fin AI Engine™ runs multi-stage safety checks (disambiguation, confidence thresholds, final safety checks)
- Claims **<1% hallucination rate**; carries **AIUC-1 audit certification** across 50+ controls

- Automatically suppresses AI answers during LLM outages
- Escalation rules via Guidance; Identity Verification gates personalized responses

The Pro add-on paywall (March 12, 2026)

Intercom moved its most valuable analytics and QA capabilities behind a paid **Pro add-on**. Existing customers get a 60-day grace period (until May 12); new customers must purchase immediately.

Six features now gated behind Pro:

1. **CX Score** — AI rates every conversation with sentiment explanation. Without Pro, CX Score is completely unavailable, including in custom reports.
2. **Topics Explorer** — Auto-groups conversations into topics/subtopics with performance data.
3. **Trends** — Weekly report identifying the largest volume shifts with plain-language cause explanations.
4. **Recommendations** — Identifies where Fin struggles, prioritizes fixes by volume/impact, generates AI content recommendations with review-edit-approve workflow.
5. **Monitors** (closed beta) — Structured QA at scale. See deep dive below.
6. **Custom Scorecards** (closed beta) — Define weighted, scorable conversation quality attributes.

What remains free: Performance dashboard (automation rate, resolution rate, involvement rate), 12 pre-built templates, chart customization, CSV export. CX Score is gated by Pro regardless of plan tier.

Paywalled surfaces: Under "Analyze" — Recommendations, Topics Explorer, Trends, Monitors, Custom Scorecards, CX Score charts. Under "Reports" — CX Score reports and chart templates.

Pro pricing (tiered by conversation volume):

Tier	Volume	Price
Base	Up to 1,000	\$99 flat
Tier 2	1,000–5,000	\$0.12/conversation
Tier 3	5,000–50,000	\$0.10/conversation
Tier 4	50,000–100,000	\$0.06/conversation
Tier 5	100,000+	Contact Sales

Intercom's worked example: 35K monthly conversations = **\$3,579/month for Pro alone**.

Monitors deep dive (closed beta, Pro required)

Monitors replace ad-hoc sampling and spreadsheet QA with a repeatable two-layer system:

- **Monitors define *what* gets reviewed** – targeting via filters (resolution state, topic, CX Score) plus natural-language "flag criteria" (e.g., "customer uses phrases such as 'cancel my subscription'"). Each monitor runs independently; multiple can flag the same conversation.
- **Custom Scorecards define *how* each conversation is evaluated** – configurable attributes (tone, accuracy, policy adherence), each scored by AI, human, or both.

Key design details:

- Attributes can be **weighted** (0–100), **marked critical** (failing = entire review score becomes 0%), and given a **pass threshold** (e.g., minimum 80%)
- Admins choose evaluation timing: "Fin is done" or "Conversation is closed"
- **Test monitor tool** validates flag criteria against real conversations before going live – shows which would be flagged, highlights mismatches
- AI scores can be overridden by human reviewers; statuses: not reviewed, reviewed, fix complete, needs action
- Out-of-the-box "Fin Quality Scorecard" provided as starting template
- Backfills recent history (Topics up to 90 days, CX Score up to 28 days)

Monitors roadmap (not yet shipped):

- Structured sampling (fixed weekly/monthly sample sizes)
- Advanced reporting (filter by monitor, scorecard, attributes, scores)
- Out-of-the-box procedure monitors (flag connector failures, execution errors, user frustration, escalation quality)
- Real-time alerts when conversations cross thresholds or fail scorecards
- Pre-deployment scoring (test changes against scorecards before going live)
- Human agent QA (monitors + scorecards for teammate conversations, not just Fin)
- Voice not yet supported

The goals gap

At Pioneer 2025, CEO Eoghan McCabe said Fin will have goals admins can set – objectives and priorities that Fin pursues while making trade-offs. **This is vision, not a shipped feature as of April 2026.** No settings UI exists.

Monitors + Custom Scorecards is the closest thing to goal-scoped evaluation, but there is **no connection between declared business outcomes and the monitors/scorecards created**. Admins

build QA criteria manually; the platform does not auto-generate from a stated goal. If Kustomer ships goals-first config where declaring "CSAT target: 90%" auto-generates both monitors and matched scorecards, that is meaningfully differentiated.

SWOT

Strengths:

- Best-in-class documentation and admin flywheel UX
- Simulations + Controlled Rollout for pre-deployment testing
- AIUC-1 certification; 40M+ conversations resolved; 66%+ resolution rates
- Monitors + Custom Scorecards (even in beta) = most structured QA framework in mid-market
- Publicly stated Goals roadmap signals narrative leadership

Weaknesses:

- Goals is vapor — not shipped
- Pro add-on paywall creates pricing backlash (10K convos/month = ~\$1,179/month *just for analytics*)
- Monitors still closed beta, not GA
- No native A/B with statistical significance
- Confidence threshold not user-configurable
- Custom Answers removal forces heavier Procedures authoring

Opportunities for Kustomer:

- Beat Intercom to shipping Goals
- Include goal-aligned analytics in the core product (exploit Pro paywall backlash directly)
- Ship goal-scoped auto-generated monitors + scorecards before Intercom connects QA to outcomes
- Ship procedure monitors, pre-deployment scoring, and human agent QA before Intercom graduates them from roadmap

Threats to Kustomer:

- Direct mid-market/upper-mid overlap
- Fin Standalone works with Zendesk/Salesforce/HubSpot/Gorgias and can land inside Kustomer accounts
- Aggressive startup program (90% off + Fin free year 1)
- Once Monitors GAs and roadmap ships, Intercom's QA surface will be very competitive

Pricing

- **\$0.99/outcome** (50-outcome minimum)
- Seats: Essential \$29, Advanced \$85, Expert \$132/seat/month (annual)
- Copilot: \$35/user/month
- Pro add-on: \$99/month base + tiered overage
- Fin Voice: custom

Stacked cost for a 20-agent mid-market team (5K outcomes, 10K conversations):

Line item	Monthly cost
Seats (Advanced × 20)	\$1,700
Copilot (× 20)	\$700
Outcomes (5,000 × \$0.99)	\$4,950
Pro add-on (10K convos)	\$1,179
Total	\$8,529

Does not include Voice, Proactive Support Plus (\$99/mo), or channel fees. Strong TCO counter-positioning opportunity for Kustomer.

Zendesk

Product capabilities

Zendesk's "Resolution Platform" (2025–2026) spans five pillars: AI Agents, Knowledge Graph, Actions & Integrations, Governance & Control, Measurement & Insights.

Product lineup:

- **AI Agents Essential** (bundled) and **AI Agents Advanced** (Ultimate.ai successor, sales-gated)
- **Agent Copilot** (\$50/agent/month)
- **Intelligent Triage** – intent, sentiment, language detection (~150 languages)
- **Zendesk QA** (Klaus) – \$35/agent/month; AutoQA scores 100% of interactions including AI agents

Relate 2025 launches: AI Agent Builder (NL procedure definition), Knowledge Builder (auto-generates KB from tickets), Action Builder, App Builder, AI Reasoning Controls (chain-of-thought visibility), AI Insights Hub.

October 2025 AI Summit: Voice AI Agents, Admin Copilot, Microsoft Copilot integration. Analytics powered by Advanced Insights (HyperArc acquisition, July 2025).

Forethought acquisition (closed early April 2026): Adds Resolution Learning Loop™ (self-improving AI – detects workflow gaps, generates procedures, tests optimizations before deployment), specialized B2B/B2C/B2E agents, autonomous workflow execution, native voice automation, "computer use" for legacy systems.

Goals-first gap

Zendesk CEO Tom Eggemeier's Relate 2025 framing – "the only metric that matters is resolution" and "outcome-based pricing" – is a **billing narrative, not a configuration UX**. No admin wizard targets a CSAT or deflection goal with auto-configuration. AI Insights Hub provides reactive recommendations; Admin Copilot surfaces guidance but doesn't prescribe from declared goals.

SWOT

Strengths:

- Massive install base (~20,000 AI customers; \$200M projected AI ARR 2025)
- Deepest partner ecosystem
- QA + WFM + Copilot + AI Agents = complete stack
- Forethought adds self-improving AI
- Strong enterprise compliance (ADPP)

Weaknesses:

- No goals-first UX
- Advanced AI Agents pricing opaque ("Talk to Sales")
- Per-resolution + per-agent stacked costs = 2–3× base plan
- No native A/B for AI agents; no historical-ticket simulation

Opportunities for Kustomer:

- Undercut on pricing transparency
- Ship goals-first before Zendesk productizes Forethought
- Position CRM-native data model as superior grounding vs. Knowledge Graph

Threats to Kustomer:

- Scale and brand authority hard to match
- Forethought's self-improving AI
- Enterprise procurement defaults favor Zendesk

Pricing

- Support Team \$19, Suite Team \$55, Suite Professional \$115, Suite Enterprise \$169/agent/month (annual)
 - Suite + Copilot bundles: Professional \$155, Enterprise \$209
 - Add-ons: Copilot \$50, QA \$35, WFM \$25, QA+WFM \$50, ADPP \$50, Contact Center \$50
 - Automated resolution overage: ~\$1.50 committed / ~\$2.00 PAYG (not publicly listed)
 - Included allowances: 5/10/15 per agent/month by tier
 - **Effective stacked cost: \$225-\$300/agent/month + per-AR fees**
-

Sierra

Product capabilities

Sierra's **Agent OS** is the most explicit "outcomes not clicks" design in the category.

Core primitives:

- **Journeys** (Agent Studio 2.0, Nov 2025) – define goals + guardrails in natural language
- **Skills** – composable action blocks with configurable determinism
- **Agent Data Platform** (Nov 2025) – cross-conversation memory scoped to strategies, audiences, outcomes, inventory
- **Ghostwriter** (March 2026) – converts SOPs, transcripts, or plain English into production-ready agents in 30+ languages

Eval rigor (category benchmark):

- Open-source: τ -bench, τ^2 -bench, τ^3 -bench
- **Simulations** (35,000+ tests/day), **Voice Sims** (dual-loop audio testing with accents, noise, emotional states)
- **Experiments** (multivariate/A/B with Control vs. Test dashboards)
- **Agent Traces** (step-by-step LLM call, tool invocation, logic branch visibility)
- **Explorer** (NL deep research across all conversations)

Governance:

- Multi-supervisor "Jiminy Cricket" model: input-filter, output-interception, subtle-action, underspecification detector
- Four-tier instruction tolerance (adversarial = zero tolerance; tone = high tolerance)
- **PCI-DSS Level 1 compliant agents** (April 2026, industry first; SiriusXM launch customer)

SWOT

Strengths:

- Genuine outcomes-first design; eval rigor unmatched
- \$635M raised, \$150M+ ARR, \$10B valuation, 40% of Fortune 50
- Customer roster: ADT, SiriusXM, Sonos, WeightWatchers, Chime, SoFi, Ramp, Cigna, Deliveroo, Rivian

Weaknesses:

- Pre-Ghostwriter: services-heavy (3–7 months, \$50K–\$200K implementation)
- Outcome-definition ambiguity causing billing disputes
- Opaque pricing; limited self-service customization

Opportunities for Kustomer:

- Win mid-market deployments where Sierra's ACV floor (~\$150K) is prohibitive
- Match governance framing with more transparent pricing
- Target Ghostwriter's mid-market push with faster time-to-value

Threats to Kustomer:

- Gold standard enterprise benchmark
- Ghostwriter specifically targets Kustomer's mid-market
- PCI L1 + voice + healthcare wins open industries Kustomer wants

Pricing

- Outcome-based, per resolution; no public list; ~\$1–\$1.50/resolution estimated
 - "Pricing blend" offered (per-conversation for routine + per-outcome for complex)
 - Annual contract minimum ~\$150K; implementation \$50K–\$200K
 - Year 1 TCO: \$200K–\$350K typical, up to \$1.5M+ enterprise
-

Decagon

Product capabilities

Core:

- **Agent Operating Procedures (AOPs)** – NL workflows compiled into validated structured logic
- **AOP Copilot** (Sept 2025) – ingests existing SOPs to auto-generate AOP skeletons
- **Decagon Duet** (Spring 2026) – AI partner that continuously rewrites AOPs from real conversation data
- **Agent Workbench** – autonomous AOP debugging
- **User Memory** (cross-session context), **Outbound Voice** (AI-initiated calls, callbacks)
- Unified Knowledge Graph ingests Zendesk, **Kustomer**, Guru, Confluence, Slack, CRM
- Git-based agent versioning with rollback

Eval framework (best-in-class alongside Sierra):

- Two-phase offline → online model: offline = LLM-as-judge, ground-truth benchmarks, simulation-based tests; online = **Experiments suite** (GA June 2025) with gradual ramp, statistical significance, configurable p-values, rollback
- **Trace View, Ask AI** (NL analytics), 3-level intent tagging, anomaly detection

Governance:

- **Watchtower** – reviews every conversation against NL rubrics with point-scoring/trend tracking (Rippling: 100% regulatory-compliance scanning)
- PII auto-redaction via Google DLP; zero-day data retention
- SOC 2 Type II, GDPR, HIPAA with BAAs

Goals-first gap

No formal "goals" object. Success metrics are defined in kickoff via Agent Product Managers, tracked through Experiments and Watchtower – but goals are implicit, not a first-class admin surface.

SWOT

Strengths:

- Eval depth (offline → online); Experiments with statistical significance; Watchtower 100% coverage
- Customer roster: Duolingo, Chime, Rippling, Notion, Hertz, Cash App, Affirm, Deutsche Telekom
- \$481M raised, \$4.5B valuation (Jan 2026)

Weaknesses:

- No goals-first UX; no public pricing; no self-service trial
- 6-week guided deployment; Agent Assist is Zendesk-only
- No unified AI+human analytics

Opportunities for Kustomer:

- Ship goals-first declarative UX; offer CRM-native unified analytics
- Compete on price transparency vs. ~\$386K median ACV

Threats to Kustomer:

- Kustomer KB integration means Decagon can land inside Kustomer accounts
- Duet + Workbench reduce switching friction; enterprise momentum building

Pricing

- No public pricing; two models: per-conversation (most customers) or per-resolution
 - Entry ~\$95K/year, median ~\$386K/year, range \$95K-\$590K+
 - Implementation included; ~30% max discount at 1M+ conversations; no per-seat fees
-

DigitalGenius

Product capabilities

Workflow-first, not goals-first. 50–60 pre-built ecommerce use cases (WISMO, returns, refunds, warranty). Portfolio: Care AI, Purchase AI, Proactive AI, Voice AI, Warranty AI (visual defect detection), Insights AI. Implementation 2–4 weeks. Integrates with Zendesk, Salesforce, Gorgias, Kustomer, Gladly, Freshdesk.

No self-serve sandbox, no A/B framework, no eval automation. Testing is CSM-led. Governance is vague with no granular guardrail configuration.

SWOT

Strengths: Deep ecommerce domain (On, Reebok, AllSaints); visual AI for warranty/returns; broad integrations.

Weaknesses: No goals-first, no evals/A/B/simulation, modest revenue (~\$962K in 2023), reported M&A activity (April 2025).

Opportunities for Kustomer: Direct displacement target in mid-market ecommerce.

Threats to Kustomer: Low — not a category-definer.

Pricing: From \$1,000/month; \$1K–\$50K implementation; ~\$39K Year 1 for 10K tickets/month.

Siena AI

Product capabilities

The **one competitor besides Sierra that explicitly markets "Set specific goals"** on its homepage.

- **CoRE engine** chains multi-step actions to achieve goals (refunds, replacements, subscription edits)
- **Automation Assistant** converts business policies into goal-driven workflows
- **Persona Studio** — channel-specific personas (Chatty, Knowledgeable, Passionate, Formal)
- **QA Agent** — evaluates 100% of conversations against plain-language criteria
- **Ask Siena** (NL analytics), **Topics Explorer** (VoC)
- Channels: email, chat, SMS, Instagram, Facebook, WhatsApp — via Gorgias, Zendesk, **Kustomer**, Dixa

SWOT

Strengths: Explicit goals-first framing; strong admin UX; Persona Studio; QA Agent; DTC case studies (HexClad, K18, Kitsch, Simple Modern).

Weaknesses: Small scale (~54 employees, ~\$5.9M revenue June 2025); many products "Coming soon"; per-conversation pricing penalizes drafts.

Opportunities for Kustomer: Match goal framing with enterprise-grade infrastructure; target DTC brands outgrowing Siena's scale limits.

Threats to Kustomer: Primarily DTC/mid-market; goals-first narrative could be amplified by others.

Pricing: No public pricing. Per-conversation model; custom quotes. No free trial.

Gorgias

Product capabilities

Rebranded January 2025 to "Conversational AI Platform for Ecommerce." Core: AI Agent (Support + Shopping Assistant), Helpdesk, Voice, SMS. AI Agent is **Shopify-only**. AI Agent 2.0 (July 2025) added multi-step actions, WhatsApp cart recovery. 3.0 teased for July 2026.

Testing is weakest in the cohort — no historical simulation, no A/B, no eval automation. Reporting has a 72-hour delay on automated-interaction stats. Confidence-based handover is not admin-tunable to numeric threshold.

SWOT

Strengths: 17,000 ecommerce merchants; deep Shopify integration; Shopping Assistant drives revenue; per-resolution pricing simplicity.

Weaknesses: No goals-first, no evals, AI resolution + helpdesk ticket **double-billing**, AI Agent Shopify-only, 72-hour reporting latency.

Opportunities for Kustomer: Upmarket displacement as Gorgias customers outgrow SMB and need evals, governance, unified analytics.

Threats to Kustomer: Lower — mostly SMB/DTC below Kustomer's core ICP.

Pricing:

- Helpdesk: Starter \$10/mo (50 tickets) → Enterprise custom (10,000+)
- AI Agent: \$0.90/resolution annual, \$1.00 monthly — but AI tickets also count as helpdesk tickets (~\$1.22–\$1.40 effective per AI resolution)

Where Kustomer wins: the strategic playbook

1. Own "Goals-first" as category vocabulary before Intercom ships it.

Ship an admin UX where step one is "What are you trying to achieve?" with four starting goal types (Deflection, CSAT, Cost Reduction, Sales Assist) that auto-generate matched procedures, simulation personas, eval rubrics, and dashboards. Name it decisively and own the term in analyst briefings before Fin 4.

2. Tie evals to declared goals — exploit Intercom's Monitors gap.

Intercom's Monitors + Custom Scorecards are the most structured QA framework in the mid-market – but they're entirely manual and disconnected from outcomes. Ship **goal-scoped evals**: declaring a CSAT goal auto-generates a simulation battery that stresses that outcome (edge cases, escalation paths, tone). This is what Intercom forces admins to build by hand.

3. Launch goal-based pricing bundles – include analytics in core.

Intercom's Pro paywall is directly exploitable. Mid-market customers now pay \$1,000–\$3,500/month just for analytics on top of outcomes, seats, and Copilot. **Kustomer Goal Packs** – bundled outcomes with analytics, evals, and dashboards included – solve the #1 buyer complaint: predictability.

4. Lead on governance transparency.

No competitor publishes a hallucination-rate dashboard. Ship: (a) public accuracy telemetry per deployment, (b) configurable numeric confidence thresholds tied to goal type, © toggleable PII/toxicity/brand-voice evaluators.

5. Sequence AIR as the goals-aligned copilot.

An AIR copilot that inherits the same declared goals as AIC – suggesting responses that optimize for the declared outcome, flagging deviations from goal-aligned rubrics – is genuinely differentiated and a cleaner architecture story than anyone else tells.

Pricing positioning

Target "**predictable bundles at Intercom TCO parity**" rather than racing to \$0.99/outcome or chasing Sierra ACVs.

- Reference: mid-market customer spends **\$8,500+/month** on Fin or **\$225–\$300/agent/month** on Zendesk's full stack
- Kustomer Goal Packs at ~\$4,500–\$8,000/month for comparable volumes – with evals, governance, and analytics included – offer **15–30% TCO savings with dramatically better predictability**
- CFO pitch: "You pay a known monthly number tied to the goal you declared; no surprise bills when demand spikes"

Conclusion

The 2026 AI-CX market has standardized on outcome-based *pricing* while missing outcome-based *configuration*. Every competitor has conceded the goals-first gap – by silence (Zendesk, Decagon, Gorgias, DigitalGenius) or by roadmapping without shipping (Intercom).

Intercom's Pro add-on paywall widens the opportunity: Kustomer can simultaneously out-feature (goal-scoped auto-generated evals) and under-price (analytics included, not paywalled) in a single product move.

The window is real but narrow — Intercom Fin 4 and Zendesk's Forethought integration both land in 2026. Moving decisively on the five recommendations above converts a defensive position into genuine category definition.

Milestone 2: Goals Based AI Setup

Goals-Based AI Setup: Bringing Goals-First vision into Kustomer AI

Admins declare an outcome. Kustomer AI configures and reports against it.

Owner	Amisha Sharma	Status	Draft for leadership review
Rolls up to	Goals-First Kustomer	Audience	Kustomer leadership
Date	May 20 2026	Stage	Definition Directional Prototype

How this fits into Goals-First Kustomer

The Goals-First vision unifies the data model, AI, rep tools, and reporting under one primitive: the business goal. The roadmap is sequenced so a foundational milestone establishes Goals as a first-class object, and every other surface becomes goal-aware on top of it.

This initiative is the AI Setup Experience expression of that vision. It picks up **Milestone 3 (Make admin setup goal-led), scoped to Kustomer AI**: admins declare the outcome they need, and the platform proposes the AI configuration (AIC procedures, AIR guidance, monitors, alerts, reporting) to deliver it. We consume the Goals schema. We don't redefine it.

The opportunity

Admins don't buy procedures. They buy outcomes: fewer escalations, faster reps, lower cost to serve, more revenue. The current product makes them assemble those outcomes themselves, touching five or more disconnected configuration areas and hoping the result reconciles in production. Most teams either over configure and lose confidence, or under configure and never see the impact they expected.

We hear the same thing from nearly every customer: Kustomer AI is powerful, but it is complex to set up. That feedback isn't just a usability complaint, it is a structural symptom. Admins are being asked to translate business intent into a long list of configuration objects, with no through line tying those objects back to the outcome they actually care about. Simplifying configuration is therefore not a side quest for this initiative, it is half of the work. If admins can't set Kustomer AI up confidently, they will never get far enough to see the outcomes it can drive.

For Kustomer, customers who don't see a measurable outcome are at risk: they renew slower, expand less, and tell a flatter story about the product. Tying configuration to outcomes is how we shorten time to value, make impact legible, and put a visible, in-product anchor under the Goals-First narrative.

What we're building

Two connected pillars, both anchored on the platform's Goals primitive.

Simpler setup, anchored on outcomes. Admins pick one or more business outcomes (reduce escalations, lift upsell, improve resolution quality, etc.) drawn from the OOTB goal taxonomy. We then ask a few targeted questions about their business, channels, and tone so we have the context to draft well. From those inputs, Kustomer generates the AI behaviors, procedures, monitors, alerts, and reports needed to support the goals. Setup goes from "configure dozens of objects" to "answer a few questions, review the plan."

Outcomes tied to recommendations. Each goal is wired to one or more monitors that track the metrics it depends on. Once live, those monitors continuously evaluate progress against target. When a metric drifts, the admin sees a specific recommendation in their workflow (tune this procedure, expand this scope, raise this guardrail), making the platform's Suggestions and Anomaly Detection visible where admins actually work.

The shift: configuration stops being a chore the admin owns end to end and becomes a conversation between the admin and Kustomer, anchored on the outcome they came to achieve.

Jobs to be done

Admin, setting up. When I'm rolling out Kustomer AI, I want to describe the outcome I'm responsible for rather than assemble the configuration myself, so I can launch quickly without becoming an AI expert.

Admin, reviewing the plan. When I review a generated AI plan, I want to see exactly what will change, where, and why, so I can trust the system before anything goes live.

Admin, operating. When my goal's metrics drift, I want Kustomer to tell me which lever to pull, not just that something is off, so I can act without retracing my whole setup.

Leader, reporting up. When my exec asks "is the AI working?", I want to point to the outcome I committed to and the trajectory we're on, so the conversation is about business impact, not feature adoption.

How success is measured

- Time to first value cut by 50%, from contract signed to first measurable AI outcome.
- 3x more admins who can articulate the outcome their AI is driving, versus today's baseline.
- +25% admin-acted recommendations per quarter post-launch.

The lagging indicator we actually care about is whether customers see and articulate the value of Kustomer AI, measured through renewal narratives, AI seat expansion, and qualitative interviews with admins and their executives.

Setup requirements

The flow surfaces as five connected pages. Each one earns its place: it either captures a decision the generator needs, or it gives the admin a moment of control before something changes in production.

1. AI Setup intake

- Confirm auto-filled company context (name, industry, website, brands, description) pulled from the account profile, with inline edit.
- Let the admin pick the audience they're configuring for: AI for Customers, AI for Reps, or both. Subsequent questions are filtered to that audience.
- Let the admin pick one or more business goals from the OOTB taxonomy, and add custom goals.
- For accounts with conversation history, surface the top topics from the last 90 days with volume, sample utterance, average handle time, and deflectability so admins know which to put under AI first.
- Show auto-suggested performance monitors derived from the chosen goals, topics, and audience. Each is editable and removable; admins can add their own up to a sensible cap.
- Accept scenario uploads (PDF, doc, image, audio, video) so we can draft procedures from real examples.
- Capture knowledge sources (help center URL, file uploads) and tone of voice.
- Show a persistent live preview of what will be generated (procedures, monitors, copilot, escalation, tone) so the admin sees their answers materialize as configuration before they commit.

2. Generated plan (Build)

- Present the generated configuration as a reviewable plan, grouped by audience and surface: AIC procedures, AIR guidance, monitors, alerts, escalation rules, tone.
- Every generated item must be traceable to the goal and inputs that produced it.
- Each item must be editable, removable, or replaceable before saving. Nothing is locked.
- Make the blast radius explicit: which automations, agents, and customer-facing surfaces this plan will touch.
- Confirm save before moving on; saving stages the plan but does not deploy it.

3. Test & evaluate

- Show a table of evaluation categories for the staged plan, with case count, audience tag, status, last run, and score.
- Let admins run any category against the staged plan and compare against the current production version.
- Provide a side panel test console for ad-hoc conversations across any channel and brand.
- Surface a readiness signal based on category scores so admins know whether the plan is safe to deploy.

4. Deploy

- Let admins scope which conversations the plan applies to via rule groups (channel, brand, customer attributes, etc.).

- For each audience, show the specific configuration that will be turned on: AIC automation behaviors, AIR copilot modes, team assignments.
- Require a deploy note and capture who deployed when. Explore how to handle collision here: eg: two admins refining an AI Agent and hoping to deploy their versions.
- Show version history with the ability to view, compare, and roll back.

5. Analyze

- Show a goal-level scoreboard: each declared goal, its target, its current trajectory, and the monitors feeding it.
- Let admins drill into conversations by version, audience, customer, channel, and date range to see exactly what the AI did.
- Surface recommendations when a metric is drifting: which procedure to tune, which guardrail to raise, which scope to expand, with one-click acceptance.
- Let admins jump directly from a recommendation back into the Build page to apply it.

Open questions

- **Pricing and packaging.** How goals and monitors will be priced and packaged is still unknown. This shapes whether goals-based setup is a default experience for every customer or a gated capability, and needs to be resolved before GA.
- Would goals be a hard requirement or can users skip goals and still configure AI?
- How much data would we store on goals? For example: These were my goals last year, vs this year and the metrics tied to them.