

PRD

Background	1
Use Cases	1
Topics in Reporting	1
Topics Configuration	1
Managing Topics	1
Topics in Data Explorer	2
Topics in Search Assistant	2
Recommendations based on Topics	2
Topics in UI	2
Open Questions	2
Competitive Analysis	2
Resources	3
Notes / Open Questions	3

Author: Amisha Sharma Víctor Peinado

Claude Prototype: <https://claude.ai/artifacts/latest/361c5ce6-6e6a-427a-ad1c-079f9ebfc879>

Background

Problem Statement: Support teams need to quickly understand what issues customers are raising and how those issues impact key outcomes like customer experience (CX) scores. Currently, Kustomer relies on manual tagging and anecdotal analysis of conversations, which is labor-intensive and inconsistent. This makes it difficult to spot recurring topics or correlate specific issue types with metrics such as satisfaction (CSAT), resolution time, or first-contact resolution.

AI Topics automatically analyze the closed conversations to uncover patterns in the questions customers ask. Using real conversation content—no setup, manual tagging, or rules required—AI groups similar conversations into meaningful topics like "Invalid credentials" or "Reset password link," which roll up into broader themes such as "Login & Account Access". Topics update continuously as new conversations are closed, helping you stay on top of emerging trends and recurring issues effortlessly.

Use Cases

Topics Report

- (P0) When I am reviewing support operations, I want to understand what topics and sub-topics are driving conversation volume (both for human reps & AI agents), so I can prioritize process improvements and resource allocation.
- (P0) When I am analyzing CX performance (both for human reps & AI agents), I want to see which topics and sub-topics correlate with poor or excellent customer satisfaction (FRT, AHT, CSAT), so I can target coaching and workflow improvements.

- (P0) Admins should be able to drill down into a list of conversations associated with a particular topic or sub-topic.
- (P1) When I am monitoring support trends, I want to see the most common questions asked by the customer per topic, so I can work on improving the knowledge resources.
- (P2) When I'm not actively working within Kustomer, I want to be alerted for trends as per my configuration.

Topics Configuration

- (P0) When I am configuring workspace settings, I want to **enable or disable Topic Analysis**, so I can comply with regulatory requirements or organizational policies.

Managing Topics

- (P0) AI Generated list of topics and sub-topics including description, created at, updated at, updated by.
- (P0) Rename topics so they match how the team describes issues.
- (P0) Merge different topics to remove duplicates. We need admin to define the new topics and description.
- (P0) Create your own new topic, and future matches will be auto-tagged into it. (Depends on the volume, but should take 24 hours to show new conversations mapping)
- (P0) Hide irrelevant or outdated topics so they do not clutter Data Explorer or Reporting.
- (P1) Save as draft to stage curation changes and publish them later.
- (P1) Custom instructions for topics
- (P1) Move descriptors between topics to put them under the right parent.

Topics in Data Explorer

- (P1) Ensure Topics and Sub-topics as a Conversation attribute are accessible in Data Explorer.

Topics in Search Assistant

- (P0) Add Topic and Sub-topics as a filter attribute on the Conversation object in Search

Topics in Team Pulse

Recommendations based on Topics

- (P1) Common Questions: per-topic view of the most frequent customer questions.
- (P1) Admins want to understand the gaps in knowledge sources for the topics that customers need support on.
- (P1) Admins want to be recommended on which topics are good candidates for automation.
- (P1) Admins want to be recommended and create evals for topics that do not have sufficient coverage.

Topics in UI

- (P0) Sub-topic should be listed in conversation attributes that are visible to the admins and can be changed by admins, if inaccurate
- (P0) Sub-topic should show up in the conversation header to give a quick view into what the conversation is about

Topics Internal Reporting

- Adoption: Track how often the topics report is viewed
- Adoption: Track how the manage topics feature is used

Open Questions

- Pricing and Packaging of Topics

Competitive Analysis

- Intercom:
<https://www.intercom.com/help/en/articles/12521409-how-to-tailor-your-ai-topics-with-topic-curation>
- Zendesk:
<https://support.zendesk.com/hc/en-us/articles/4550640560538-Automatically-detecting-customer-intent-sentiment-and-language>

Resources

- Detailed PRD:  Core Topic Reporting (LDX PRD)
-  Topics Analysis

Notes / Open Questions

- M1 Scope:
 - Generating Topics (Delta-2)
 - Topics Clustering (Category): Combine similar topics (Delta-2)
 - Managing Topics (Delta-2) Blair Li
 - Show Topics in UI (Delta-2) Jon Chin : Reps need a quick way to understand what the conversation is about. Admins can change the name of the topic.
 - Top 3 Similar Conversations in UI Jon Chin : Reps need a quick way to review similar conversations to take inspiration to resolve the current conversation.
 - Admins can change the topic of a conversation (Delta-2) Blair Li
- Topics Analysis Report (OOTB) (Alpha-2)
- Topics in Data Explorer (Alpha-2)
- Topics Filtering in Existing OOTB Reports (Alpha-2)


- Topics Filters in New Reports (Alpha-2)
- M2 Scope:
 - Centralize the suggestions from AI
 - AI Agent Suggestions based on Topics
 - Coaching Suggestions on Topics
- We can show similar conversations.

AI-generated topics and descriptors in KAlpha

AI-generated topics and descriptors in KAlpha

Every conversation marked as Done will be automatically analyzed in real time to generate AI annotations: up to 5 granular descriptors and 1 broad topic. What do we actually mean with *topic* and *descriptors*? See these FAQ: [AUT-689 Conversation Topics](#)

These annotations will show up in the conversation view, right under the conversation name.



They are stored at the conversation level. For now, they're static and cannot be edited. They can be manually regenerated, by reopening and closing the conversation again.

```
{
  "phase": "active",
  "annotations": {
    "topic": "product issues",
    "descriptors": [
      "defective water bottle",
      "order ORD-993776",
      "product failure",
      "warranty inquiry",
      "return policy"
    ],
    "createdAt": "2026-03-11T20:38:50.553Z"
  }
},
```

Notice these AI-generated annotations will always be in English, regardless of the language of the conversation. That's fine because we will use their embeddings (which are multilingual) to actually capture the meaning of the conversation.

在数字包容性至关重要的时代，满足需求的能力 对于全球观众来说，这不仅是一种优势，而且是一种必要。这世界以丰富的语言进行交流，其中包括西班牙语...

我对客服团队非常失望。真倒霉，我竟然要经历这一切。以后再也不会从这家公司订购任何东西了。

Topic: **language options**

Descriptors: **language support** **interface options** **dissatisfaction**

2/13/2026

✓ Done

Сразу два образа знаменитого режиссёра отразил на одной картине (на илл.) его друг-художник

Сразу два образа знаменитого режиссёра отразил на одной картине (на илл.) его друг-художник.

Topic: **account settings change**

Descriptors: **account settings** **change notification** **issue unresolved**

2/13/2026

✓ Done

The instructions we use right now to generate the topics and descriptors are the following. In the near future, we should support custom instructions.

The actual customizations are still TBD, but as of today, I see the **RED** sections as fixed system prompts and the **GREEN** ones as likely to be overridden with custom instructions.

None

Instructions

You are a documentalist and analytics expert. Your task is to analyze the conversation summary and extract meaningful, searchable labels for analytics, automation planning, and knowledge management.

What to focus on (highest priority)

- The customer's need/intent (what they are trying to do): e.g., cancel subscription, reset password, change delivery address
- The expressed issue (what's wrong / what failed): e.g., login error, double charge, missing package
- Products, services, plans, features, and brands mentioned (use the exact proper name when available): e.g., Walmart+, iPhone 16 Pro, Air Jordans
- Policies/processes only when they are the subject of the customer's need: e.g., return policy, identity verification, refund eligibility

What to ignore (do NOT label)

- Tone/sentiment/emotion and conversational dynamics: frustration, angry, rude, polite, appreciative, calm, confusion (unless it directly names a concrete issue)
- Agent performance or meta talk: helpful agent, long wait, escalation, transfer, call quality, "asked for more info," "agent said..."
- Conversation metadata: channel (chat/phone/email), timestamps, greetings/closings, pleasantries
- Pure internal workflow unless it maps to a customer-facing need (e.g., "internal note," "ticket created," "handoff")

Descriptor rules

- Generate up to 5 atomic, granular descriptors (1-3 words each when possible).
- Prefer nouns/noun-phrases that are actionable and content-based.
- Prioritize: (1) product/service name(s), (2) core issue, (3) customer intent, (4) key policy/process (if relevant).
- Avoid redundant support-domain words:
 - Do NOT include: customer, support, customer support, inquiry, issue, problem, ticket
 - Write "return policy" not "customer inquiry about returns"
- Avoid agent-only steps unless explicitly central to the customer's need:
 - "identity verification" is OK if it blocks access/resolution
 - "information verification process" -> "identity verification"
- If multiple specific products/features are named, include the most central one(s) only (don't list everything).
- Prefer canonical phrasing:
 - "password reset" (not "resetting password")
 - "order tracking" (not "track order")
 - "subscription cancellation" (not "cancel subscription")

Topic rule

- Generate one broader topic (2-3 words) that best encompasses the descriptors (not a sentence; no tone/meta).

Examples (good)

- prior authorization, prescription billing, insurance coverage -> topic: pharmacy billing
- return policy, refund, shipping fee -> topic: returns & refunds
- identity verification, account access, password reset -> topic: account access
- Air Jordans, delivery delay, order tracking -> topic: order delivery

Examples (bad -> fix)

- "customer frustration" -> (omit) OR replace with the concrete issue causing it (e.g., "billing error")
- "agent helpfulness" -> (omit)
- "long wait time" -> (omit)

Output format

Respond with ONLY a JSON object. No markdown, no explanation.

```
{"topic": "<2-3 word topic>", "descriptors": ["<d1>", "<d2>", ...]}
```

Blair-Topic project scope

Topic project scope

1. Topic analysis (data visualization)
 - a. Topic in standard reporting - **Zainab**
 - i. Topic report includes both human and AI agent
 1. Filter by human only, AI only or both
 - b. Topic in data explorer - **Steve**
 - c. **Topic as a part of insight (benefit: contextual)** - **Zainab**
 - i. Customer timeline level
 - ii. Company timeline level
 - iii. Each convo level
2. Topic management/config - **Zainab**
 - a. setting/platform
 - b. Kustomer AI/setting
3. Topic as a new filter or new attribute - **Steve**
 - a. Topic in search
 - b. Topic in search assistant
 - c. Topic filter/new column in team pulse
 - d. Topic attribute in custom report build modal

CREATE NEW CHART

1

2

3

4

Select DataVisualizationSelect AnalysisCustomize Chart

Datasets

Choose a dataset below to see its available metrics.

Search datasets...

Standard

Conversation

Conversation-level performance insights like response time.

Customer

Customer-level engagement insights.

Message

Insights on individual messages.

Note

Insights on conversation notes and attachments.

Task Views (Handle Time)

Insights on tasks.

Work Item

Routing analysis.

Work Session

Capacity and availability analysis.

Metric library

Choose a metric to select your analysis.

No metrics to show yet

Choose a dataset to show its available metrics

Cancel

Next

Definitions for Conversation Topics Elements

Definitions for Conversation Topics Elements

1. Descriptors

Definition: Atomic, factual attributes extracted from a resolved conversation that capture a single specific aspect of the interaction.

Characteristics:

- **Granularity:** Each descriptor isolates one distinct feature (not compound concepts)
- **Objectivity:** Based on observable content, linguistic signal, not inferred sentiment
- **Consistency:** Same situation should produce same descriptor across conversations (we're somehow forcing this in the prompt, see [📄 Topics Analysis PRD](#))

Categories of descriptors:

Category	Purpose	Examples
End-user's intent and need	What the customer wants to achieve	"request refund", "track order", "cancel subscription"
Core issue	The problem or situation being discussed	"billing error", "product defect", "delivery delay"
Product/Service	What's being discussed	"mobile app", "premium plan", "golden crocks"
Policies, processes	Observable emotional/situational context	"return policy", "refund eligibility", "verification"
Resolution	Outcome or action taken	"escalated", "resolved", "pending response"

2. Broad Topic

Definition: A single categorical label that represents the primary domain or area of concern for the conversation, under which all descriptors logically fit.

Characteristics:

- **Hierarchical:** Sits one level above descriptors as a parent category
- **Mutually exclusive:** Each conversation belongs to exactly one broad topic

- **Business-aligned:** ideally, it should map to organizational departments, processes, or service areas. Our current generic instructions will cover 80% of the customer support use-cases, they will never be perfect without client's feedback/guidance or their own taxonomy.

Relationship to descriptors: The broad topic answers "What area of the business does this conversation concern?" while descriptors answer "What specifically happened within that area?"

Examples:

Broad Topic	Possible Descriptors
"Billing"	"refund request", "payment failed", "invoice dispute", "monthly subscription"
"Delivery"	"tracking issue", "late shipment", "wrong address", "youth lacrosse stick"
"Account Access"	"password reset", "locked account", "2FA issue", "ubisoft+"

3. Cluster Label

Definition: A synthesized, human-readable name for a group of conversations that share similar descriptor patterns, generated after clustering is complete.

Characteristics:

- **Emergent:** Derived from analyzing what grouped conversations have in common
- **Distinctive:** Differentiates this cluster from other clusters
- **Actionable:** Helps stakeholders quickly understand what the cluster represents

Relationship to broad topics: While a broad topic is assigned per-conversation before clustering, a cluster label is assigned per-group after clustering. Clusters may span multiple broad topics or subdivide a single topic, depending on the sample of conversations used for the clustering job.

Examples:

Cluster contents	Cluster label
Conversations about failed payments + angry customers + premium plans	"Premium Payment Failures"

Cluster contents	Cluster label
Conversations about tracking + delivery delay + first-time contact	"Delivery Delay Inquiries"

How the three elements relate

None

Conversation → Descriptors (multiple) → Clustering

↓

Broad Topic (one)

↓

↓

↓

Cluster Label (post-hoc)

- **Descriptors** are the building blocks (what happened)
- **Broad Topic** is the domain (where it belongs)
- **Cluster Label** is the pattern name (what group it's part of)

Topics: Pricing Proposal

Topics Analysis — Pricing Proposal

Kustomer AI · May 2026 · Prepared by: Amisha Sharma

Summary

Topics Analysis automatically extracts (closed conversations only) and clusters conversation topics across all Kustomer orgs using an unsupervised AI pipeline (LLM embeddings → UMAP/HDBSCAN). No training data, no human labeling. The recommendation is to offer it **free** as a core part of the Kustomer AI platform.

Cost Breakdown

Data sourced from Victor Peinado's production monitoring (May 2026). Current real-time costs reflect **all orgs** in prod. Batch costs today reflect only the pilot orgs (prod1: ~6% of full volume; prod2: ~23%). Full-rollout batch estimates use Victor's methodology: real-time cost ÷ 2 (batch API is ~50% cheaper).

Environment	Daily Cost — Real-Time (All Orgs)	Daily Cost — Batch GA (Est.)	Monthly Est. — Batch GA
Prod1	~\$112	~\$56	~\$1,680
Prod2	~\$5.15	~\$2.58	~\$77
Combined	~\$117	~\$58–59	~\$1,750–1,760

Key drivers:

- Batch processing (hourly jobs, 3–20 min runtime) is ~50% cheaper than real-time via OpenAI Batch API pricing
 - Once real-time is fully disabled (next milestone), steady-state cost lands at ~\$1,750–1,760/month across all production environments
 - Current batch figures are partial — full-volume estimates are extrapolated, pending Victor's dashboard confirmation
-

Recommendation: Include Free in Kustomer AI Platform

Absorb the ~\$1,750/month as a strategic platform cost. Do not gate behind a paywall.

Why:

1. **AI-native positioning** — Competitors (Intercom Fin Topics Explorer, Zendesk, Salesforce) are all investing here. A paywall weakens Kustomer's out-of-the-box AI value story.
 2. **Platform feature, not an add-on** — Topics feeds downstream value: reporting, routing, Copilot context, search enrichment. Fragmenting it breaks the value chain.
 3. **Cost is manageable and declining** — ~\$1,750/month is reasonable relative to engineering investment made. Batch migration already cut costs in half; further optimization is possible.
 4. **Adoption velocity matters** — Alpha orgs (KAlpha, Turo, Zwift, Vuori, Whisker, Catawiki) show strong signal. Frictionless rollout accelerates the feedback loop that improves topic quality.
-

Decisions Requested

1. Approve ~\$1,750/month Topics Analysis cost absorbed into Kustomer AI platform COGS
 2. Confirm Topics Analysis ships free to all orgs on Kustomer AI plans — no usage gate
 3. Establish a quarterly cost review cadence with a defined threshold for revisiting pricing
-

Cost data: Victor Peinado, #project-topics · Linear KUS-10508 · Note: Batch API pricing in dashboards pending Victor's confirmation — final numbers to be verified before GA.